

# The Transport Layer

---

- ◆ Having examined the TCP/IP Protocol Reference the focus now shifts to the TCP layer.
  - This layer is the heart of the whole protocol hierarchy.
- ◆ It sits below the Application Layer providing a ‘Transport’ service to the applications that wish to communicate across a Network.

# The Transport Layer

---

- ◆ TCP provides a reliable, cost-effective end-to-end data transport service *independently* of the physical network(s).
- ◆ It is important to realise that a *service* is very different to a *protocol*, although they are frequently confused.

# Services

---

- ◆ Each layer of the reference model provides a set of functionality to the layer immediately above:
  - This set of functionality is known as “The Services”.
- ◆ The layer below is known as the “Service Provider” and the layer above is known as the “Service User”
- ◆ The Services are accessed through the interface between the layers.

# Protocols

---

- ◆ Protocols on the other hand are how the Services are implemented.
- ◆ Typically a Protocol specifies a framing structure which will include a number of fields containing Control data:
  - Generically this framing structure is known as a Protocol Data Unit (PDU)
  - Examples include Data Link Frame, IP Datagram etc.

# Protocols

---

- ◆ The Protocol will also typically specify a procedure for interpreting and responding to the Control data within the PDU.
- ◆ For instance, if a service provides for reliable transfer of data then there must be some means of tracking and recovering from data loss.

# Protocols

---

- ◆ In this instance part of the PDU Control field will include numbering:
  - e.g. byte numbers, frame numbers etc.
- ◆ The Protocol will also specify how data in the Control Fields are to be interpreted and responded to if necessary.
  - e.g. for missing frame return a REJ message.

# The TCP Transport Service Offering

---

- ◆ The TCP Transport Service has the following characteristics:
  - **Connection Orientation:** Before two applications entities can communicate they must first establish a connection.
  - **Point-To-Point Communication:** Each TCP connection has exactly two endpoints.
  - **Complete Reliability:** TCP guarantees that the data will be delivered exactly as sent i.e. no data missing or out of sequence
  - **Full Duplex Communication:** A TCP connection allows data to flow in either direction
    - TCP buffers outgoing and incoming data
    - This allows applications to continue executing other code whilst the data is being transferred

# The TCP Transport Service Offering

---

- ***Stream Interface***: The source application sends a *continuous* sequence of octets across a connection
  - The data is passed *en bloc* to TCP for delivery
  - TCP does not guarantee to deliver the data in the same size pieces that it was transferred by the source application.
- ***Reliable Connection Startup***: TCP both applications to agree to any new connection
- ***Graceful Connection Shutdown***: Either can request a connection to be shut down
  - TCP guarantees to deliver all the data reliably before closing the connection



# The Transport Service

---

- ◆ The TCP transport service is offered to a *user process* that exists within the application layer:
  - This *user process* is considered a “Transport Service User”,
  - The service is typically offered through a set of **primitives** across the interface between the layers,
  - Calls to these primitives cause the transport service provider to perform some action.

# The Transport Entity

---

- ◆ To aid future discussions the TCP software within the transport layer that implements the service will be referred to as the *Transport Entity*:
  - This is the “Transport Service Provider”.

# The Transport Entity

---

- ◆ The *transport entity* can be located in any number of places including:
  - Within the operating system (OS) kernel,
  - As a separate user process,
  - Within a library package bound to the network application,
  - On the network interface card.

# The Transport Entity

---

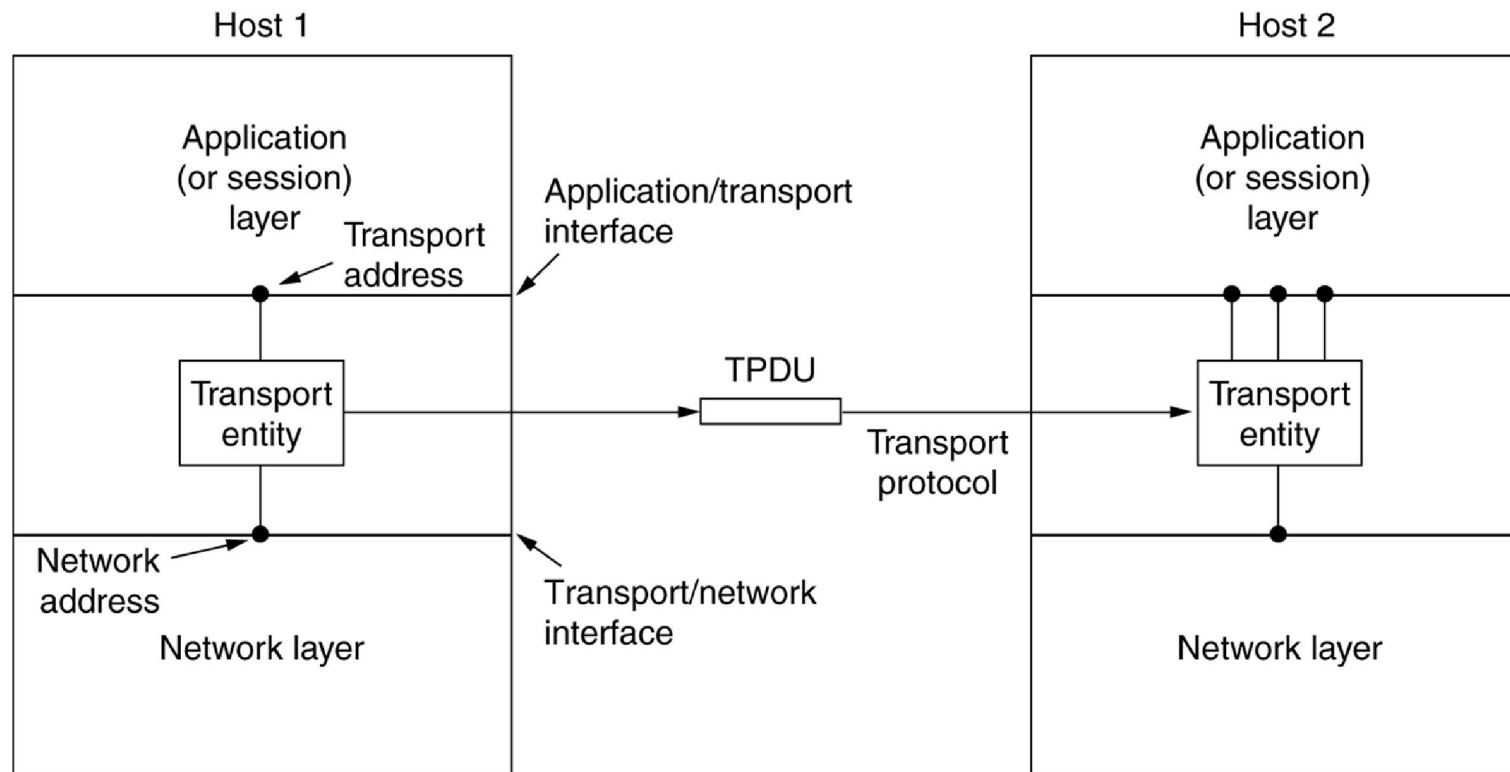
- ◆ If the *protocol stack* is located within the OS the **primitives** are implemented as *system calls*:
  - ◆ These calls turn control of the machine over to the OS to send and receive the necessary PDUs.
- ◆ The next diagram shows the Transport Entity in context.
- ◆ As can be seen it shows the Transport Layer sitting between the Application and Network layers.

# The Transport Entity

---

- ◆ It has a connection to each of the layers above and below:
  - It is the “Service Provider” to the Application Layer and,
  - A “Service User” of the Network Layer.
- ◆ The TPDU represents the framing structure that is exchanged between *Peer Entities*:
  - The PDU does NOT move horizontally between the Transport layers,
  - Instead it moves up-and-down through the Protocol Stack.

# The Network, Transport and Application Layers



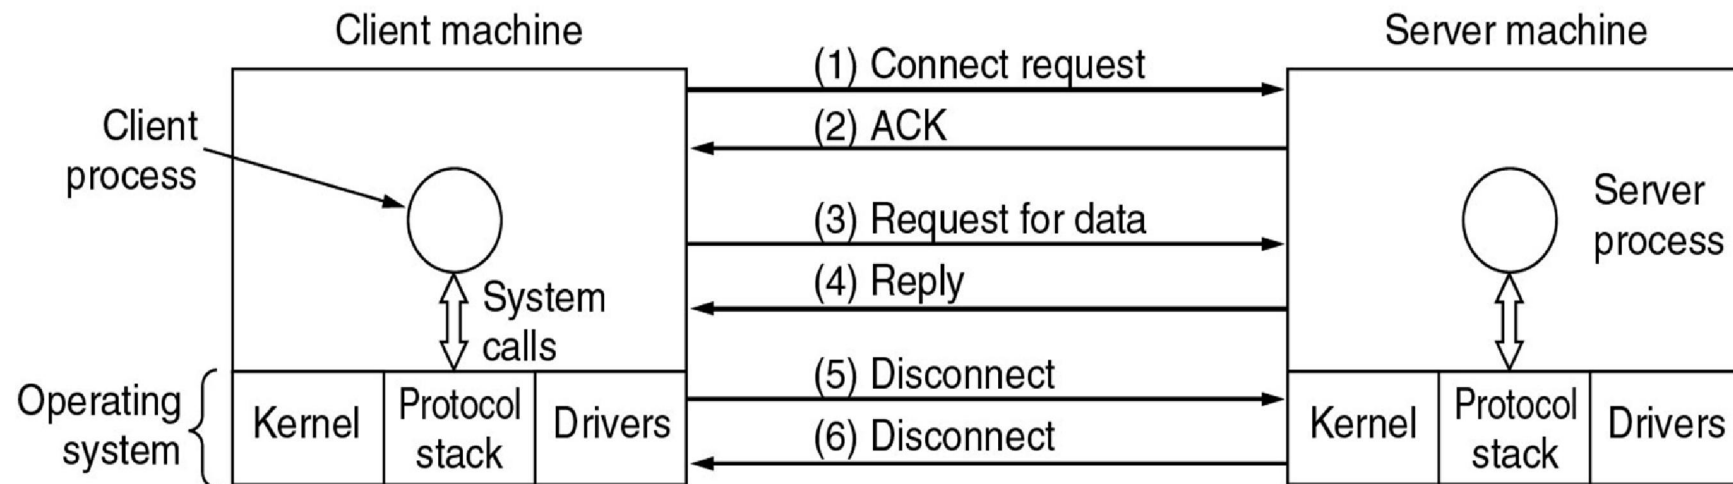
# Transport Service Primitives

---

- ◆ What do these primitives look like?
- ◆ Consider a ***generic*** transport interface as shown in the next diagram:
  - Here can be seen a list of Primitives.
  - These primitives allow application programs to establish, use and release connections.
- ◆ Typically there is a very precise sequence of calls to these primitives:
  - The exact sequence differs for *clients* and *servers*.

# A set of basic Transport Service primitives

Primitive	Packet sent	Meaning
LISTEN	(none)	Block until some process tries to connect
CONNECT	CONNECTION REQ.	Actively attempt to establish a connection
SEND	DATA	Send information
RECEIVE	(none)	Block until a DATA packet arrives
DISCONNECT	DISCONNECTION REQ.	This side wants to release the connection





# Managing connections

---

- ◆ Before explaining the sequence of primitive calls it is important to understand the *Phases of Communication*.
- ◆ Recall from previous discussions on HDLC that there are generally three phases of communication associated with a connection-oriented services:
  - Connection Establishment,
  - Data Transfer, and,
  - Connection Release.
- ◆ During each phase a variety of PDUs are exchanged between the *client* and *server*:
  - Each PDU contains a message for the “other side”.

# Connection interactions per phase

---

## ◆ *Connection Establishment Phase:*

- The server firstly executes a LISTEN primitive:
  - The server's *transport entity* responds to this primitive call by blocking the server until a client request arrives,
- The client's then executes a CONNECT call:
  - This causes a CONNECTION REQUEST TPDU to be sent to the Server,
- The server transport entity unblocks the server and returns a CONNECTION ACCEPTED TPDU to the Client,
- The client transport entity then unblocks the client and the connection is deemed **established**.

# Connection interactions per phase

---

## ◆ *Data Transfer Phase:*

- With an active connection data can now be exchanged between the client and server using the SEND and RECEIVE primitives,
- Each side must take turns using (blocking) RECEIVE and SEND.

## ◆ *Release Phase:*

- Either the client or the server can call a DISCONNECT primitive:
  - ◆ This sends a DISCONNECT TPDU to the remote transport entity
  - ◆ Upon arrival, the connection is deemed released.

# Berkeley Sockets

---

- ◆ The basic transport primitives discussed above are not standardised.
- ◆ Instead most OS designers have adopted the *socket primitives*
  - These originated from the Berkeley University of California's UNIX OS which contained the *TCP/IP* suite of internetworking protocols
- ◆ Consequently the socket API has become the *de facto* standard for interfacing to TCP/IP

# Socket Communication and UNIX

---

- ◆ Coming from a UNIX background *sockets* use many concepts found in UNIX
- ◆ An application communicates through a *socket* in the same way that it transfers data to or from a file
- ◆ For this UNIX uses an *open-read-write-close* paradigm
  - An application calls:
    - *open* to prepare a file for reading/writing
    - *read* or *write* to retrieve or send data to/from the file
    - *close* to finish using the file
- ◆ When *open* is first called a *descriptor* is returned
  - All future interaction with the file requires this *descriptor*
  - Socket communication also uses this descriptor approach

# Socket Communication And UNIX

---

- ◆ To use the TCP/IP protocols to communicate an application must request the OS to create a **socket**:
  - This is an abstract concept which will be explained in detail later.
- ◆ The OS returns a descriptor that *identifies* the socket:
  - For now just consider it as a *file handle*.
- ◆ As with File I/O this descriptor must be used in all future interactions with the socket
- ◆ The following slides lists the Socket API primitives and illustrates an example of how these primitives are used by a client and server application

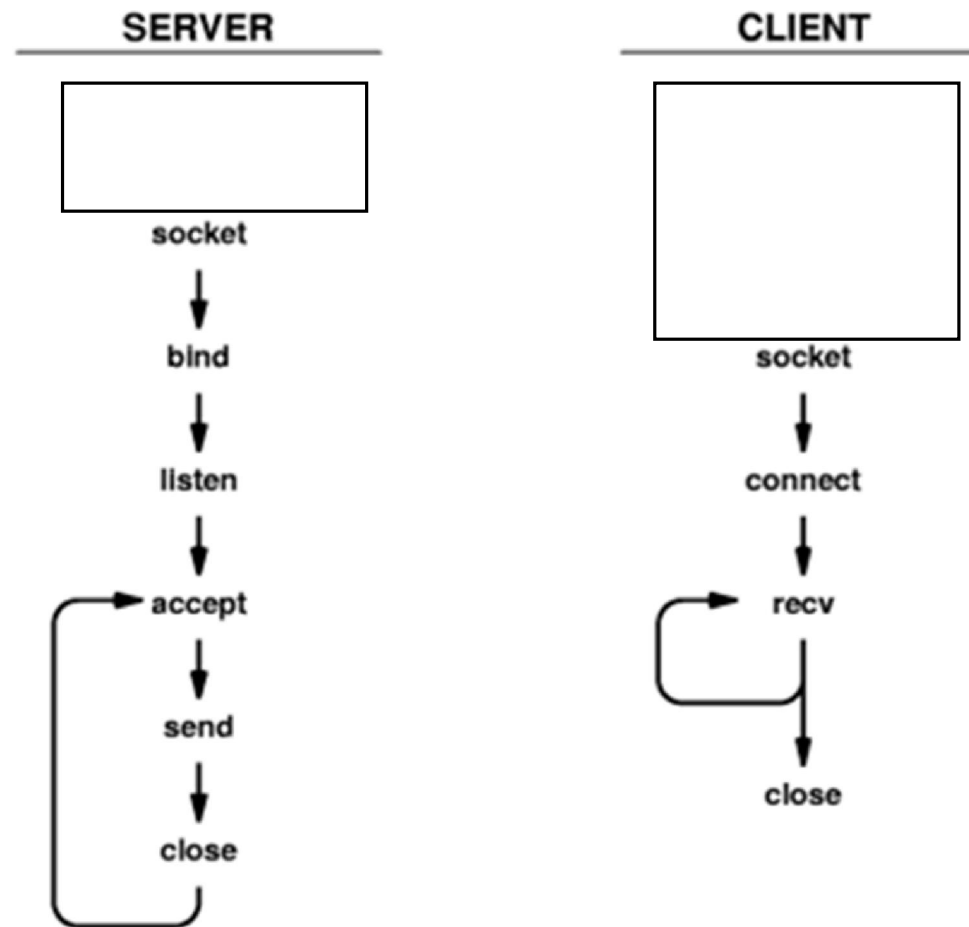
# The Socket Transport Primitives

---

Primitive	Meaning
SOCKET	Create a new communication end point
BIND	Attach a local address to a socket
LISTEN	Announce willingness to accept connections; give queue size
ACCEPT	Block the caller until a connection attempt arrives
CONNECT	Actively attempt to establish a connection
WRITE/SEND	Send some data over the connection
READ/RECEIVE	Receive some data from the connection
CLOSE	Release the connection

## An Example Client-Server Interaction Using Sockets

---





# The Socket Primitives - explained

---

- ◆ The following primitives are executed by *servers*:
  - SOCKET: This primitive creates a new *end point* within the Transport Entity:
    - Table space is allocated within the transport entity,
    - A file descriptor is returned which is used in all future calls
  - BIND: This primitive binds a socket to a network address:
    - This allows remote clients to connect to it

# The Socket Primitives - explained

---

- LISTEN: This primitive allocates a queuing space within the transport entity for incoming call requests.
- ACCEPT: This primitive blocks the server waiting for an incoming connection:
  - Upon receipt of a connection request the *transport entity* creates a new socket identical to the original one and returns a file descriptor to the server.
  - The server *forks off* a new process or *service thread* to handle the connection on the new socket
  - The server also continues to wait for more connections on the original socket

# The Socket Primitives - explained

---

- ◆ The following primitives are executed by *clients*:
  - SOCKET: As before
  - CONNECT: This primitive blocks the client and actively starts the connection process
  
- ◆ The following primitives are executed by *clients* and *servers*:
  - SEND and RECV: These primitives are used to transmit and receive data over the full-duplex connection
  - CLOSE: This primitive releases the transport connection

# Sockets and Socket Libraries

---

- ◆ In most systems the *socket* functions are part of the OS.
- ◆ Some systems, however, require a *socket library* to provide the interface to the *transport entity*:
  - These operate differently to a *native* socket API,
  - The code for the library socket procedures are linked into the application program and resides in its address space,
  - Calls to a socket library pass control to the library routine as opposed to the OS.
- ◆ Both implementations provide the same semantics from a programmer's perspective
- ◆ Applications using either implementation can be **ported** to other computer systems.

# Example use of Sockets

---

- ◆ The next slide shows an example Client application using the Sockets API.
- ◆ It is a simple Daytime Client which connects to a Daytime Server for the current date and time on the Server Host.
- ◆ This application is written in 'C'.
- ◆ Lines 24, 41, 44 and 57 show four of the socket primitives being called.
- ◆ This programme will be explained in class and you will have a chance to compile and run it against a pre-compiled server.

# An Example Client Using Sockets

---

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <string.h>
4  #include <unistd.h>
5  #include <sys/types.h>
6  #include <sys/socket.h>
7  #include <netinet/in.h>
8  #include <arpa/inet.h>
9  #include "Practical.h"
10
11  int main(int argc, char *argv[]) {
12      char recvbuffer[BUFSIZE]; // I/O buffer
13      int numBytes = 0;
14
15      if (argc < 3) // Test for correct number of arguments
16          DieWithUserMessage("Parameter(s)",
17                              "<Server Address> <Server Port>");
18
19      char *servIP = argv[1]; // First arg: server IP address (dotted quad)
20
21      in_port_t servPort = atoi(argv[2]);
22
23      // Create a reliable, stream socket using TCP
24      int sock = socket(AF_INET, SOCK_STREAM, IPPROTO_TCP);
25      if (sock < 0)
26          DieWithSystemMessage("socket() failed");
27
28      // Construct the server address structure
29      struct sockaddr_in servAddr; // Server address
30      memset(&servAddr, 0, sizeof(servAddr)); // Zero out structure
31      servAddr.sin_family = AF_INET; // IPv4 address family
32      // Convert address
33      int rtnVal = inet_pton(AF_INET, servIP, &servAddr.sin_addr.s_addr);
```

# An Example Client Using Sockets

---

```
34     if (rtnVal == 0)
35         DieWithUserMessage("inet_pton() failed", "invalid address string");
36     else if (rtnVal < 0)
37         DieWithSystemMessage("inet_pton() failed");
38     servAddr.sin_port = htons(servPort);    // Server port
39
40     // Establish the connection to the echo server
41     if (connect(sock, (struct sockaddr *) &servAddr, sizeof(servAddr)) < 0)
42         DieWithSystemMessage("connect() failed");
43
44     while ((numBytes = recv(sock, recvbuffer, BUFSIZE - 1, 0)) > 0) {
45         recvbuffer[numBytes] = '\0';    // Terminate the string!
46         fputs(recvbuffer, stdout);    // Print the echo buffer
47         /* Receive up to the buffer size (minus 1 to leave space for
48          * a null terminator) bytes from the sender */
49     }
50     if (numBytes < 0)
51         DieWithSystemMessage("recv() failed");
52         // else if (numBytes == 0)
53         // DieWithUserMessage("recv()", "connection closed prematurely");
54
55     fputc('\n', stdout); // Print a final linefeed
56
57     close(sock);
58     exit(0);
59 }
60
```