

Addressing Information

- ◆ Very often there is a need to exchange addressing information between the Application layer and the TCP layer.
- ◆ Both the client and server need to do this after the socket has been created.
 - The client needs to pass the contact details for the server before the Connect Primitive is called.
 - The server needs to inform the TCP of the address that it wants to listen on. This is used by the Bind Primitive.
- ◆ Occasional there is a need to pass information in the reverse direction i.e. TCP to Application layer. This will be looked at another time.
- ◆ All addressing information regardless of direction must be of the correct byte order (discussed shortly) and can only be passed by *reference* through a standardised address structure.
- ◆ This structure is specified by the Sockets API and is discussed in the next slide.

Socket Address Structures

```
struct in_addr
{
    in_addr_t s_addr;          /* 32-bit IPv4 address network byte
ordered */
};
```

```
struct sockaddr_in
{
    uint8_t sin_len;           /* length of structure (16) */
    sa_family_t sin_family;    /* AF_INET */
    in_port_t sin_port;        /* 16-bit TCP or UDP port number
network byte ordered */
    struct in_addr sin_addr;    /* 32-bit IPv4 address
network byte ordered */
    char sin_zero[8]; /* unused */
};
```

- ◆ Only concerned with three members in the structure:
 - *sin_family*, *sin_addr*, and *sin-port*.
 - The *sin_zero* member pads the structure to at least 16 bytes in size

Socket Address Structures

- ◆ Socket address structures are of local significance
 - i.e. they are not communicated between different hosts
- ◆ Socket address structures are always passed by reference
- ◆ However *socket functions* are designed to deal with socket address structures from many supported protocol families
- ◆ These functions do not understand the the generic pointer type (*void **)
- ◆ Consequently a *generic* socket address structure of the following form was created:

The *Generic* Socket Address Structure

```
struct sockaddr
{
    uint8_t      sa_len;
    sa_family_t  sa_family;    /* address family: AF_XXX value */
    char         sa_data[14];  /* protocol-specific address */
};
```

- ◆ The socket functions are defined as taking a pointer to this generic socket address structure
- ◆ For example in the *bind* function:

```
int bind(listenfd, (SA *) &servaddr, sizeof(servaddr));
```

where, SA is defined as *struct sockaddr* and
serveraddr was declared as (*struct sockaddr_in*).₄

Byte-order

- ◆ Two ways to store a 16-bit integer that is made up of 2 bytes:
 - Store the low-order byte at the starting address, known as *little-endian* byte order
 - Store the high-order byte at the starting address, known as *big-endian* byte order
- ◆ The byte ordering used by a given system is known as the *host byte order*. Unfortunately there is no standard between these two byte orderings

Byte-order

- ◆ Client and server applications will typically that extend across host systems that use both formats.
- ◆ Consequently programmers of networked applications must deal with the *byte-ordering* differences as follows:
 - TCP uses 16-bit port number and a 32-bit IPv4 addresses.
 - Both end-protocol stacks must agree on the order of these bytes
 - TCP/IP uses *big-endian* byte ordering and is known as *network byte order*
- ◆ Certain fields within the socket address structures must be converted from *host byte order* to *network byte order*

Byte Ordering and Manipulation Functions

- ◆ Depending on the level of conversion there are two sets of functions that can be used;
 - Byte Ordering functions are the simplest in that they deal with string-to-numeric-to-string conversion
 - Byte Manipulation functions are more complex in that they deal with more complicated string manipulation i.e. from dotted-decimal notation-to-numeric-to-dotted-decimal notation
- ◆ There are four byte ordering functions:
 - htons() – converts host 16-bit value to network byte order
 - htonl() – converts host 32-bit value to network byte order
 - ntohs() – converts 16-bit network value to host byte order
 - ntohl() – converts 32-bit network value to host byte order

Byte Manipulation Functions

◆ **inet_pton**

- This function takes an ASCII string (*presentation*) that represents the destination address (in dotted-decimal notation) and converts it to a binary value (*numeric*) for inputting to a socket address structure (i.e. *network byte order*)

◆ **inet_ntop**

- This function does the reverse conversion, i.e. from a *numeric* binary value to an ASCII string representation (*presentation*) i.e. *dotted-decimal notation*

Byte Manipulation Functions

- ◆ Both functions work with IPv4 and IPv6 addresses

```
#include <arpa/inet.h>
```

```
int inet_pton ( int family, const char *strptr, void *addrptr);
```

- Returns: 1 if OK, 0 if input not a valid presentation format, -1 on error
- convert the string pointed to by *strptr*, storing the binary result through the pointer *addrptr*

```
const char *inet_ntop(int family, const void *addrptr, char *strptr, size_t len) ;
```

- Returns: pointer to result if OK, NULL on error

- ◆ The *family* argument for both functions is AF_INET as we are only concerned with IPv4 addresses.
- ◆ The *len* argument is the size of the destination buffer and is passed to prevent the function from overflowing the buffer

The *accept* Function – Capturing Client addresses

- ◆ Accept is called by a server to return the next connection from the connection queue:
 - If the queue is empty, the process is put to sleep

```
#include <sys/socket.h>
```

```
int accept ( int sockfd, struct sockaddr *cliaddr, socklen_t *addrlen ) ;
```

- Returns: nonnegative Descriptor if OK, -1 on error

The *accept* Function – Capturing Client addresses

- ◆ *accept* returns up to three values:
 - An integer return code that is either a new socket descriptor or an error indication,
 - The *protocol* address of the client process (through the *cliaddr* pointer)
 - The size of this address (through the *addrlen* pointer)

The *accept* Function – Capturing Client addresses

- ◆ The *cliaddr* and *addrlen* arguments are used to return the *protocol* address of the client process:
 - Before the call to *accept* is made:
 - **addrlen* is set to the size of the client address structure (*cliaddr*),
 - On return this integer value contains the actual number of bytes stored by the kernel in the socket address structure.
- ◆ If we are not interested in the address of the client the last two arguments are set to **NULL** pointers.

The *accept* Function – an example

```
3      int main(int argc, char **argv)
5      {
        .
7      socklen_t clntAddrLen; // new variable to hold length of address structure
8      struct sockaddr_in servaddr, cliaddr; //new address structure
9      char clntName [INET_ADDRSTRLEN]; //buffer to hold the client address
        .
19     clntAddrLen = sizeof(cliaddr); //determines length of Client Address Structure
20     connfd = Accept(listenfd, (SA *) &cliaddr, &clntAddrLen); // call to accept
21     printf("connection from %s, port %d\n", inet_ntop(AF_INET,
        &cliaddr.sin_addr, clntName, sizeof(clntName)),
        ntohs(cliaddr.sin_port)); //print out client address
        .
```