

Name and Address Conversions

- ◆ All examples examined so far have intentionally used *IP addresses*:
 - These IP addresses are used by clients connecting to servers,
 - However, end-users rarely know of, let alone use, the IP address(es) of the servers they wish to communicate with,
 - Instead, we use **hostnames**.
- ◆ Remote host computers are normally known by **human-readable names**:
 - Using *names* instead of numbers is easier for humans,
 - It also allows the *numeric IP address* for a remote host to change without changing the *hostname*.
- ◆ However, client applications connecting to remote server applications still require an IP address:
 - Some method is required to map hostnames to IP addresses.

Name and Address Conversions

- ◆ A *Name service* is used to map between *hostnames* and *IP addresses* (amongst other things):
 - The process of mapping a *hostname* to a numeric quantity such as an IP address or Port Number is called **resolution**,
 - When an IP address for a particular hostname is obtained from a *name service* the hostname is said to be **resolved**.
- ◆ Two primary name service sources are:
 - The *Domain Name System* (DNS). This is a distributed name service requiring the use of the DNS protocol. and,
 - Local configuration databases which are operating-system specific.
- ◆ Fortunately from a programming perspective the details of the name service are hidden:
 - Programmers only need to know how to ask for a name to be **resolved**.

Resolvers and Name Servers

- ◆ Organizations often run one or more *name servers* (DNS server).
- ◆ Applications contact a DNS server by calling functions in a library known as a ***resolver***:
 - The *resolver* code is contained in a system library and is link-edited into the application during the *build* process,
- ◆ Calls to the *resolver* code are made using functions such as ***getaddrinfo ()*** and ***getnameinfo ()***
 - The former maps a *hostname* into its IP address, and the latter does the reverse mapping

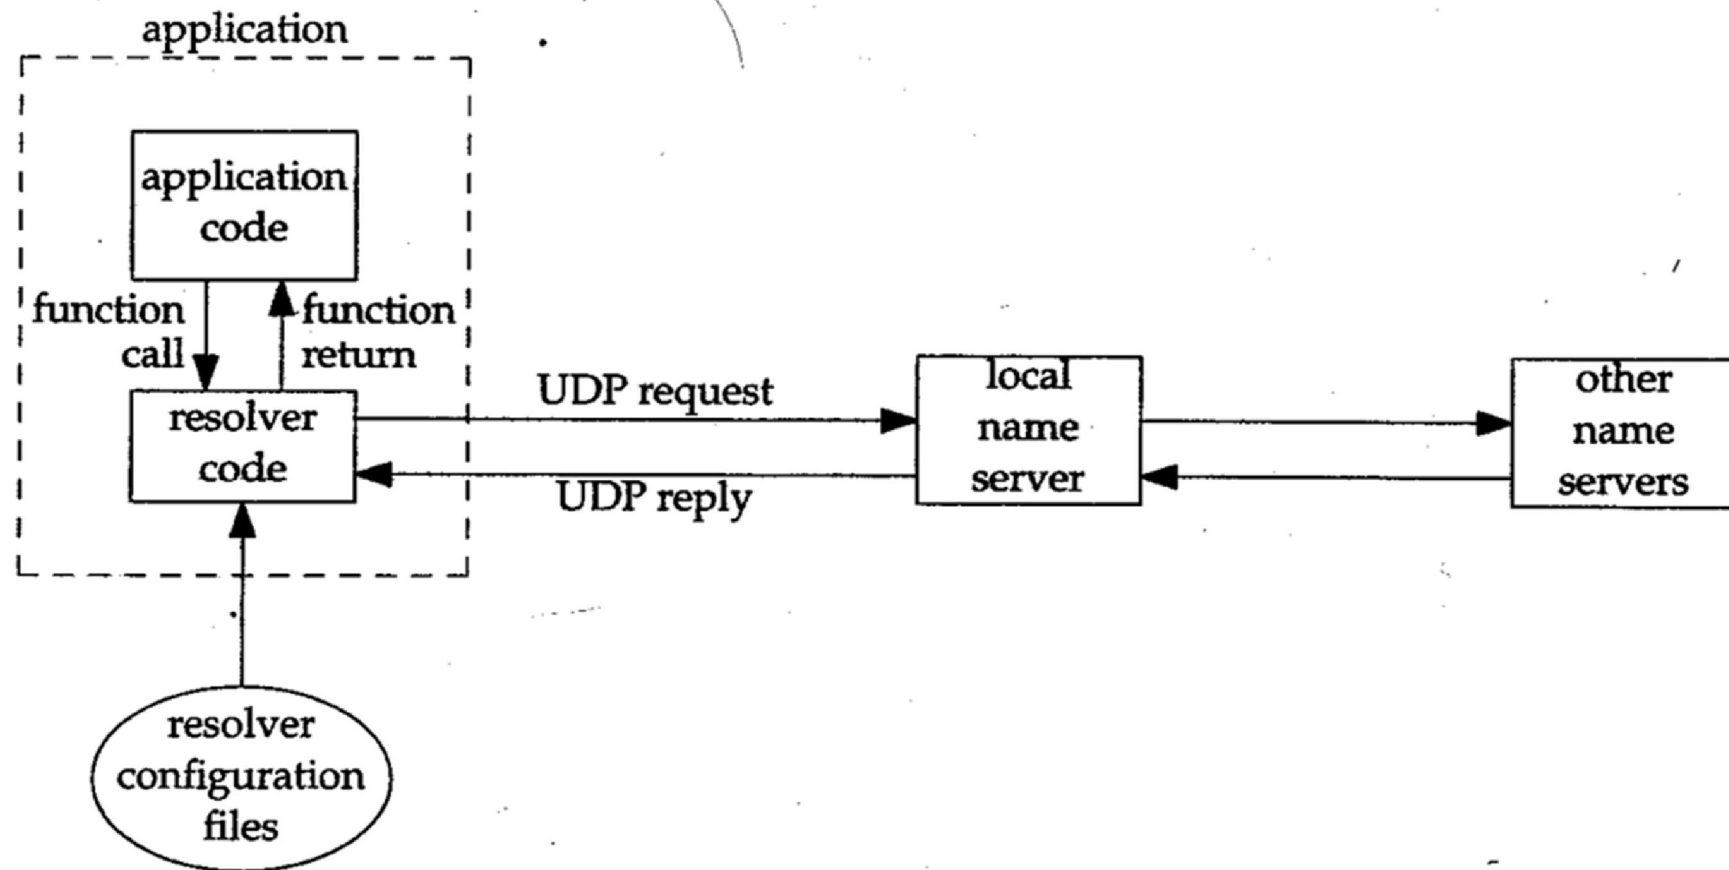
Resolvers and Name Servers

- ◆ The resolver code refers to a *configuration file* to determine the location of the name server(s)
 - The file `/etc/resolv.conf` normally contains the IP address of the local name servers
- ◆ The resolver sends the query to the local name server:
 - If necessary the local server will query another name server

Name and Address Conversions

- ◆ Entries in the DNS are known as *resource records* (RRs):
 - *A* records map hostnames to a 32-bit IPv4 address,
 - *AAAA* records map hostnames to a 128-bit IPv6 address,
 - *PTR* records map *IP* addresses into *hostnames*,
 - *MX* records specifies a host to act as a *mail exchanger*,
 - *CNAME* records map common services, such as *ftp* and *www* to the actual host providing the service
 - ◆ Example www.dit.ie has the *canonical* name *remus.dit.ie*.
- ◆ The RRs that we are interested in is the ***A record***.

Resolvers and Name Servers



The *getaddrinfo()* function

- ◆ ***getaddrinfo*** () performs a query for an *A* record:

```
int getaddrinfo (const char *hostStr, const char *serviceStr,  
const struct addrinfo *hints, struct addrinfo **results);
```

Returns: NULL if OK and a non-error code if unsuccessful.

e.g. **getaddrinfo** ("www.google.com", 0, NULL, &addrList);

- ◆ ***hostStr*** points to a null-terminated character string representing a **host name** such as *aisling*, or, a fully qualified domain name (FQDN) such as: *aisling.student.dit.ie*
- ◆ ***serviceStr*** will be ignored for the moment
- ◆ ***hints*** describes the kind of information to be returned
- ◆ ***results*** is the location of a *struct addrinfo* pointer which points to a linked list containing the results i.e. the protocol addresses

The *getaddrinfo()* function

- ◆ Each entry in the linked list is held in an *addrinfo structure* which is declared as follows:

```
struct addrinfo {  
    int          ai_flags;    // Flags to control information resolution  
    int          ai_family;   // Family: AF_INET, AF_UNSPEC etc.  
    int          ai_socktype; // Socket type: SOCK_STREAM,  
                               SOCK_DGRAM  
    int          ai_protocol; // Protocol: 0 (default) or IPPROTO_XXX  
    socklen_t    ai_addrlen;  // Length of socket address ai_addr  
    struct sockaddr *ai_addr;  // Socket address for socket  
    char         *ai_canonname; // Canonical name  
    struct addrinfo *ai_next;  // Next addrinfo in linked list  
};
```


The *getaddrinfo()* function

- ◆ The **ai_addr** field contains a *sockaddr* structure of the appropriate type, populated with (numeric) address and port information.
- ◆ In the case of TCP/IP the appropriate type is a *sockaddr_in* structure which has the following members:

```
struct sockaddr_in
{
    uint8_t          sin_len;          // length of structure (16)
    sa_family_t      sin_family;       // AF_INET
    in_port_t        sin_port;         // 16-bit TCP or UDP port number
                                         network byte ordered
    struct in_addr    sin_addr;         // 32-bit IPv4 address
                                         network byte ordered
    char             sin_zero[8];      // unused
};
```

The *getaddrinfo()* function

- ◆ The linked list of results returned by ***getaddrinfo()*** must be deallocated:
 - This requires the use of the auxiliary function, ***freeaddrinfo()***
 - Given a pointer to the head of the linked list it frees all the storage allocated for the list. Failure to call this method can result in a memory leak
- ◆ The following example program (GetAddrInfo.c) illustrates the use of ***getaddrinfo()*** and ***freeaddrinfo()***:
 - The program takes two command-line parameters, a hostname and a service name (or port number), and prints the IP address(es)
./GetAddrInfo www.dit.ie http

A sample program using *getaddrinfo()*

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <string.h>
4  #include <netdb.h>
5  #include "Practical.h"
6
7  int main(int argc, char *argv[]) {
8
9      if (argc != 3) // Test for correct number of arguments
10         DieWithUserMessage("Parameter(s)", "<Address/Name> <Port/Service>");
11
12     char *addrString = argv[1]; // Server address/name
13     char *portString = argv[2]; // Server port/service
14
15     // Tell the system what kind(s) of address info we want
16     struct addrinfo addrCriteria; // Criteria for address match
17     memset(&addrCriteria, 0, sizeof(addrCriteria)); // Zero out structure
18     addrCriteria.ai_family = AF_INET; // The TCP/IP address family
19     addrCriteria.ai_socktype = SOCK_STREAM; // Only stream sockets
20     addrCriteria.ai_protocol = IPPROTO_TCP; // Only TCP protocol
21
22     // Get address(es) associated with the specified name/service
23     struct addrinfo *addrList; // Holder for list of addresses returned
24     // Modify servAddr contents to reference linked list of addresses
25     int rtnVal = getaddrinfo(addrString, portString, &addrCriteria, &addrList);
26     if (rtnVal != 0)
27         DieWithUserMessage("getaddrinfo() failed", gai_strerror(rtnVal));
28
29     // Display returned addresses
30     for (struct addrinfo *addr = addrList; addr != NULL; addr = addr->ai_next) {
31         PrintSocketAddress(addr->ai_addr, stdout);
32         fputc('\n', stdout);
33     }
34
35     freeaddrinfo(addrList); // Free addrinfo allocated in getaddrinfo()
36
37     exit(0);
38 }
```

The *getaddrinfo()* function

- ◆ Line 16 is a *struct* that is used to restrict the type of information to be returned:
 - In this case we are only interested in TCP/IP addresses
 - The location of this *struct* is passed as the third argument (**hints**) to ***getaddrinfo()***
- ◆ Line 25 is the call to ***getaddrinfo()***
- ◆ Lines 30 iterates over each node of the linked list
- ◆ Line 31 prints the IP address and Port number from the *ai_addr* member of the current linked list node:
 - ***ai_addr*** points to a *struct* of type ***sockaddr_in***
 - The addresses are then taken from the ***sin_addr*** and ***sin_port*** members of this ***sockaddr_in*** structure