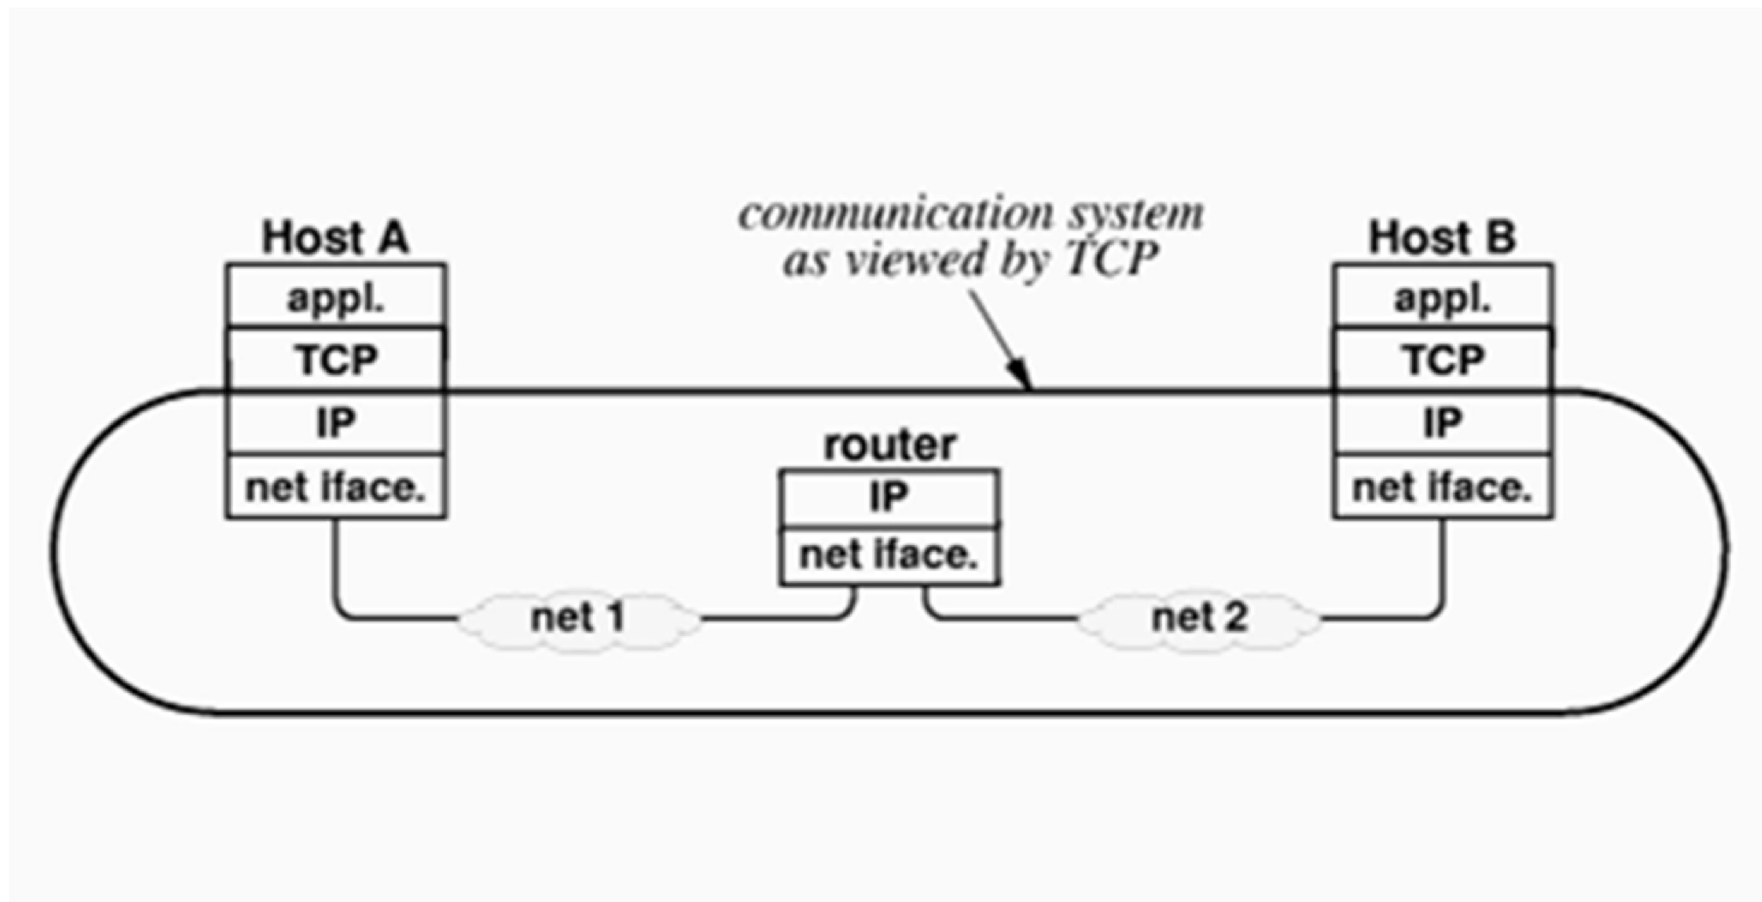


# Where TCP and IP Operate

---



# The complete TCP Transport Service offering

---

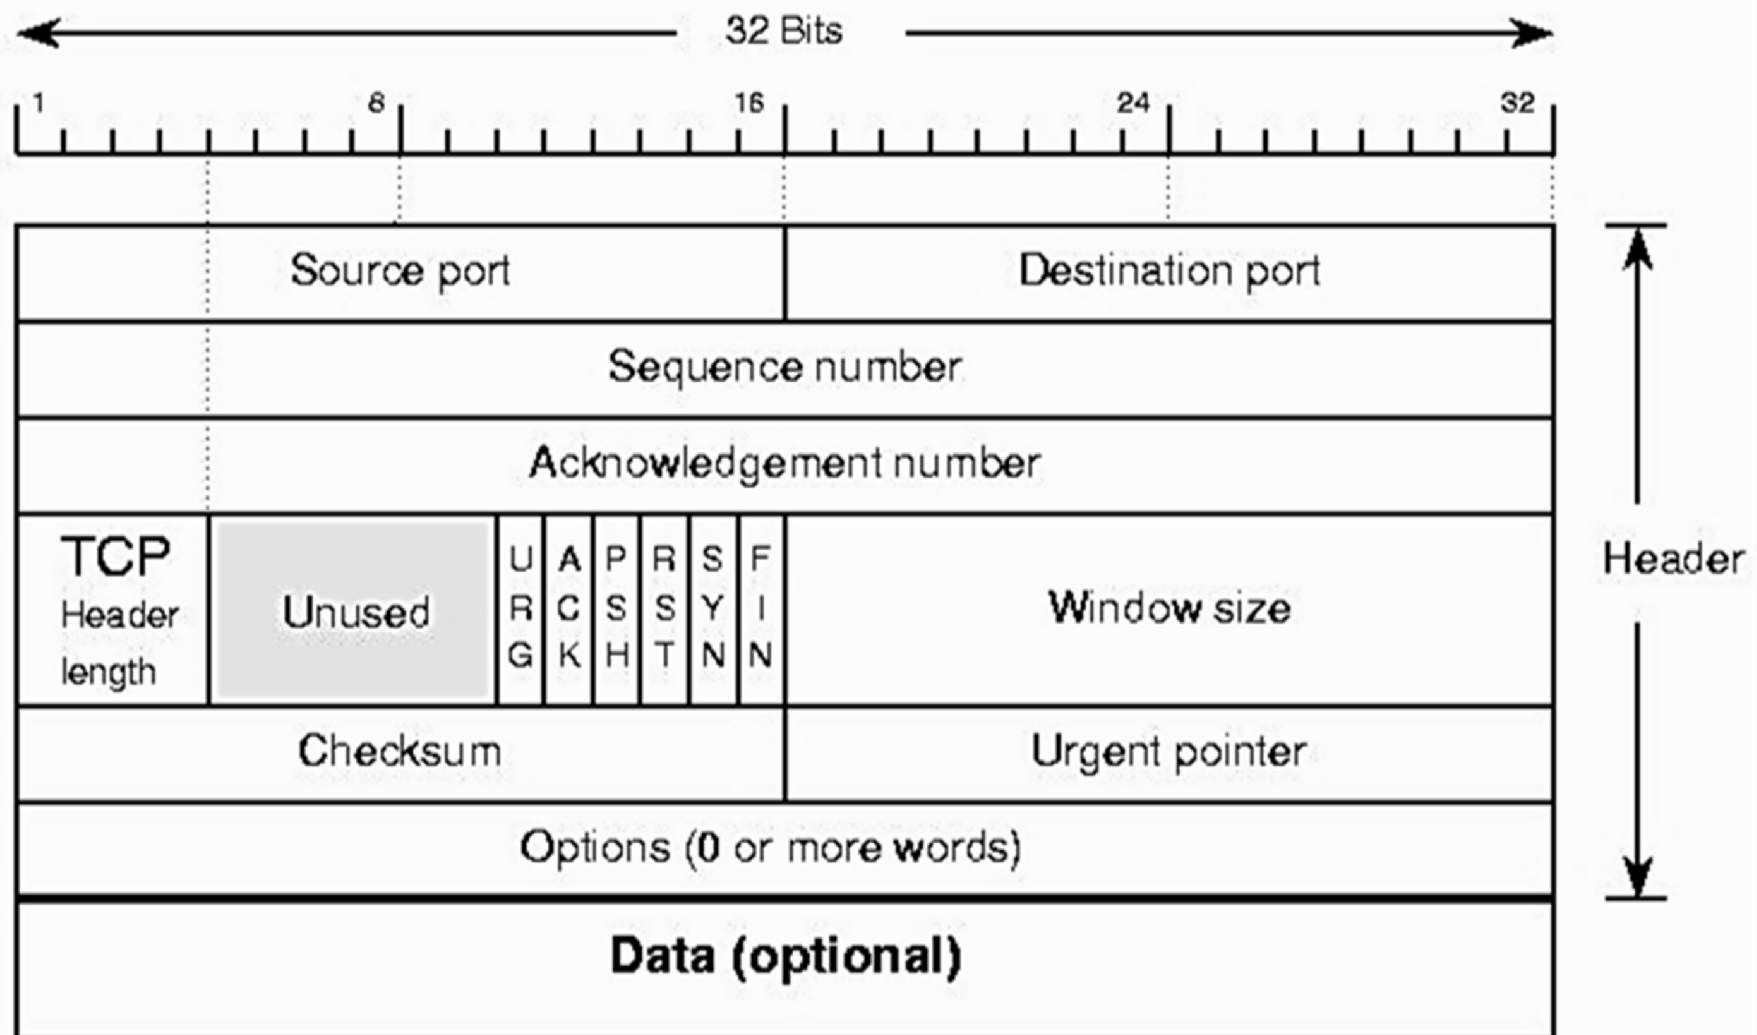
- ◆ The TCP Transport Service has the following characteristics:
  - **Connection Orientation:** Already discussed above
  - **Point-To-Point Communication:** Each TCP connection has exactly two endpoints.
  - **Complete Reliability:** TCP guarantees that the data will be delivered exactly as sent i.e. no data missing or out of sequence
  - **Full Duplex Communication:** A TCP connection allows data to flow in either direction
    - TCP buffers outgoing and incoming data
    - This allows applications to continue executing other code whilst the data is being transferred

# The complete TCP Transport Service offering

---

- ***Stream Interface***: The source application sends a *continuous* sequence of octets across a connection
  - The data is passed *en bloc* to TCP for delivery
  - TCP does not guarantee to deliver the data in the same size pieces that it was transferred by the source application.
- ***Reliable Connection Startup***: TCP both applications to agree to any new connection
- ***Graceful Connection Shutdown***: Either can request a connection to be shut down
  - TCP guarantees to deliver all the data reliably before closing the connection

# TCP Segment Header Format



# Opening a TCP connection: **The Three-Way Handshake**

---

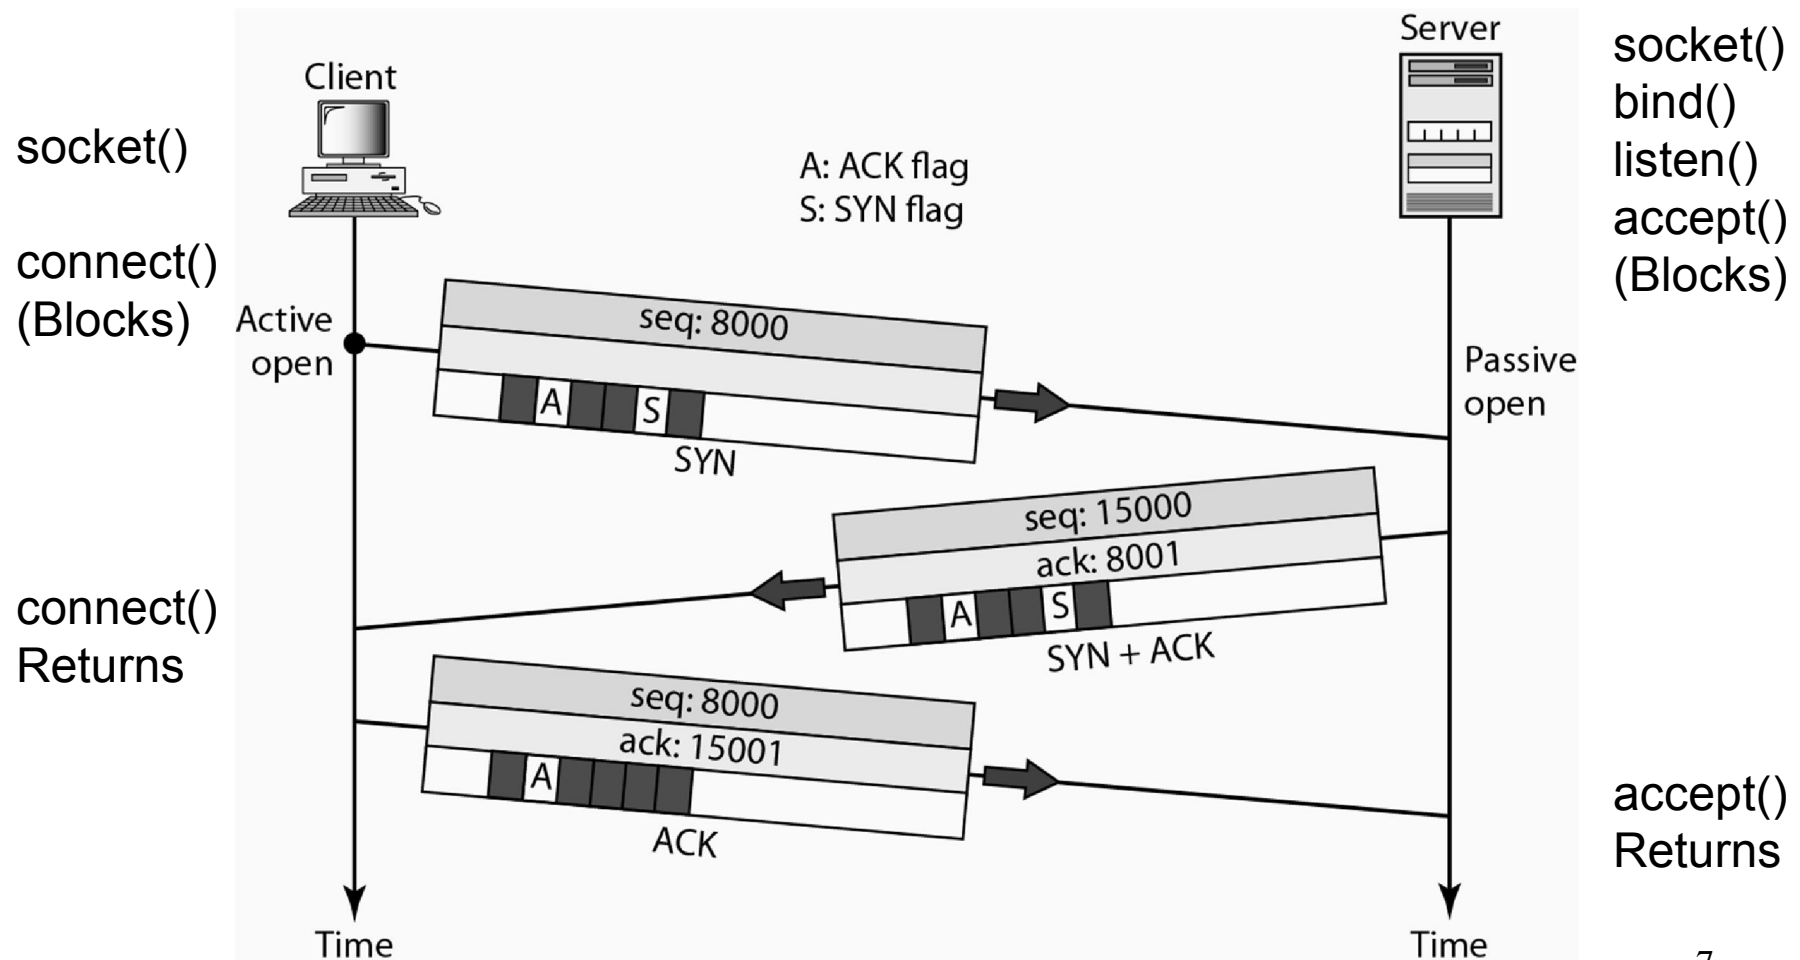
- ◆ Opening a TCP connection:
  - A server calling *socket*, *bind*, and *listen* is performing a *passive* open
  - A client calling *connect* is performing an *active* open
- ◆ The call to *connect* causes the client TCP entity to send a SYN (*synchronize*) segment:
  - This segment contains the client's initial sequence number for its data
- ◆ The server TCP entity must acknowledge receiving a SYN segment and it must send its own SYN segment:
  - This contains the initial sequence number for its data (note: the full duplex operation)
  - In addition the server's SYN contains an ACK of the client's SYN message

## The Three-Way Handshake – Contd.

---

- ◆ The client TCP entity must also acknowledge receipt of the server's SYN segment:
  - As there are three segments used in the opening sequence this procedure is called a *three-way handshake*
- ◆ **Note on sequence numbers:**
  - Sequence numbers contained in a TCP acknowledgement segment (ACK) is the next expected sequence number
  - SYN messages occupy 1 byte of the sequence number space
  - ACKs on their own do not consume a sequence number
  - Refer to the example connection start-up sequence on the next slide

## Example *Opening* Three Way Handshake



# Closing a TCP connection

---

- ◆ Closing a TCP Connection:
  - An application calling *close* is performing an *active* close
  - The other end of connection is said to be performing a *passive* close
- ◆ The call to ***close()*** causes the local TCP entity to send a FIN segment:
  - This implies that the application is finished sending data
  - Either application, the client or the server, can call ***close()***



# Closing a TCP connection

---

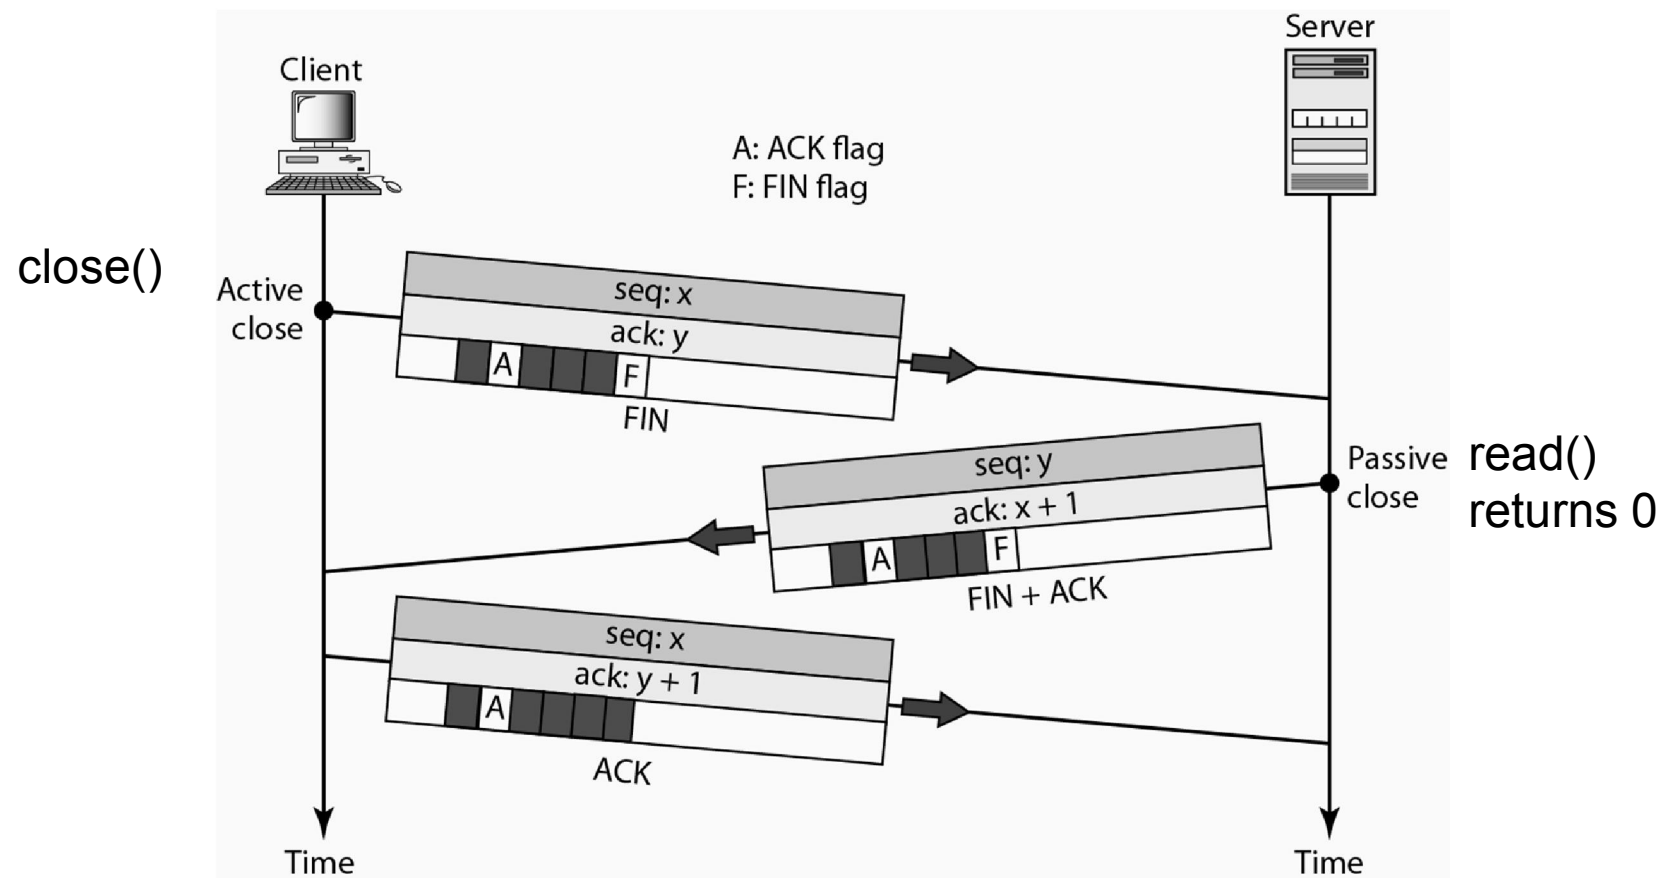
- ◆ The FIN segment is acknowledged by the receiving TCP entity:
  - This FIN message is passed to the application as an *end-of-file* and is queued after any remaining data
  - Receipt of a FIN means no more data will arrive on the connection
- ◆ An application that receives an *end-of-file* can choose to:
  - Leave the local socket open for longer in order to return data to the remote app. This is known as a *half-close*, **OR**,
  - Close its local socket by calling *close()* resulting in its local TCP entity sending a FIN. This FIN segment must be acknowledged by the remote TCP entity
- ◆ This choice results in either a *three-way* or *four-way* handshaking sequence for closing the connection

# Closing a TCP connection

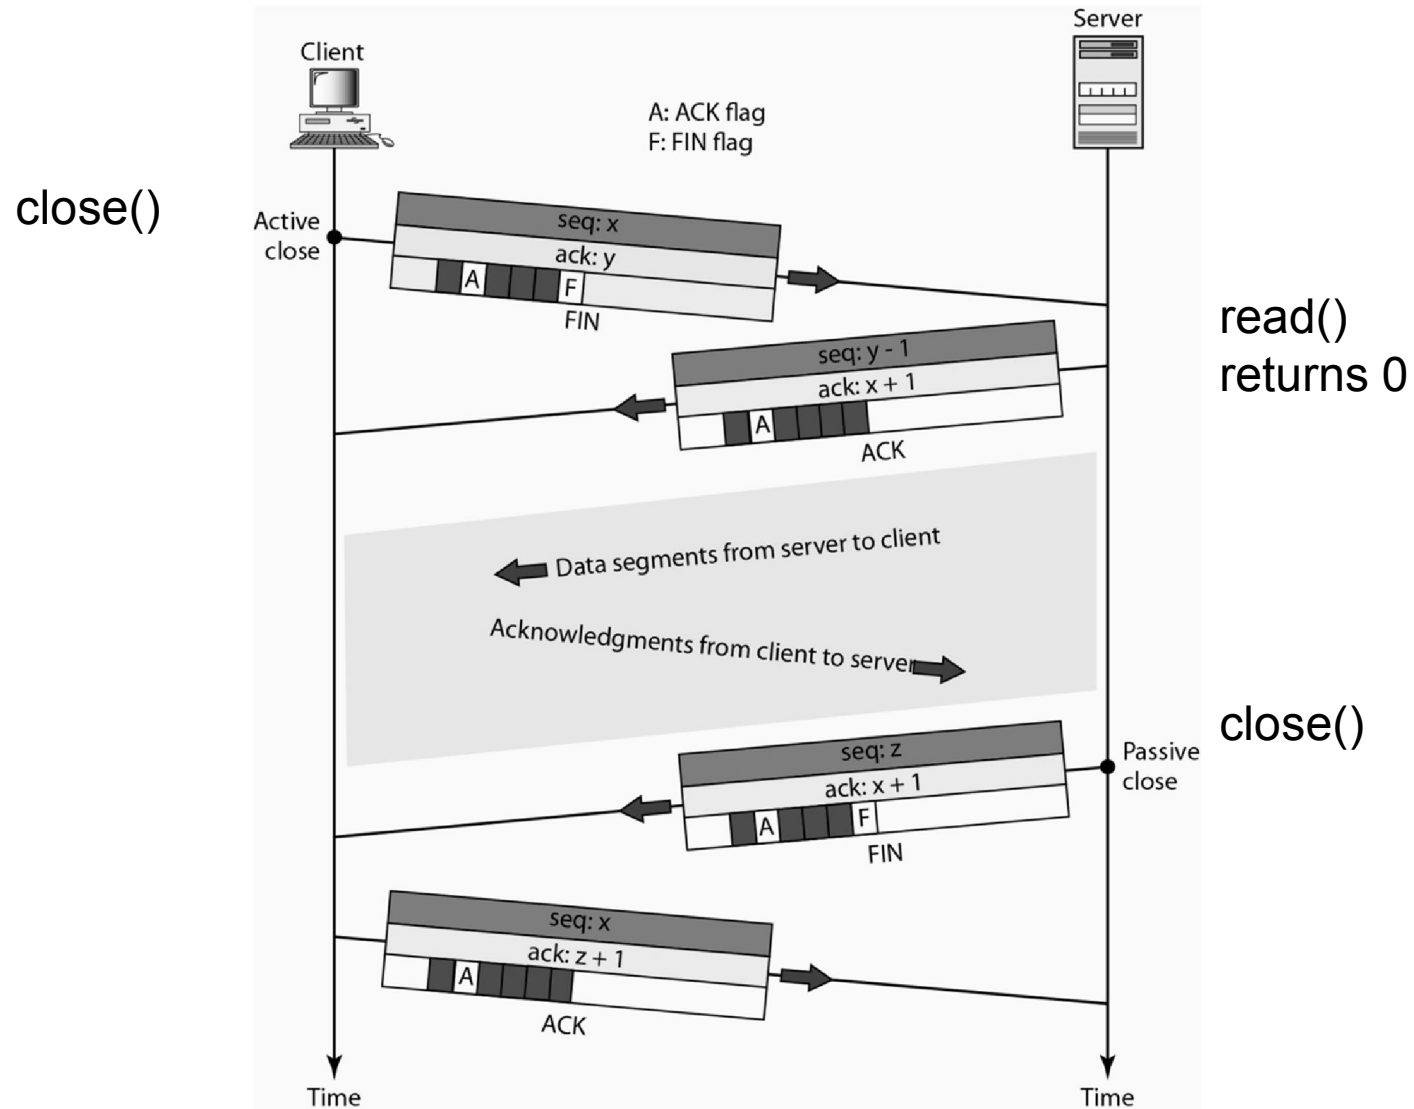
---

- ◆ Closing a connection requires a FIN and an ACK in each direction i.e. four segments are normally required:
  - However the *active* FIN is normally sent with data and the *passive* FIN can be combined with the ACK into a single segment i.e. a three-way handshake
- ◆ Connections can also be closed by terminating the application i.e. terminating the associated Unix process:
  - This causes all open socket descriptors to close
  - This results in a FIN segment being sent on any open TCP connections
- ◆ **Note on sequence numbers:**
  - Just like SYN segments, FIN segment also occupy 1 byte of the sequence number space
  - ACKs on their own do not consume a sequence number
  - Refer to the example connection termination sequence on the next slide

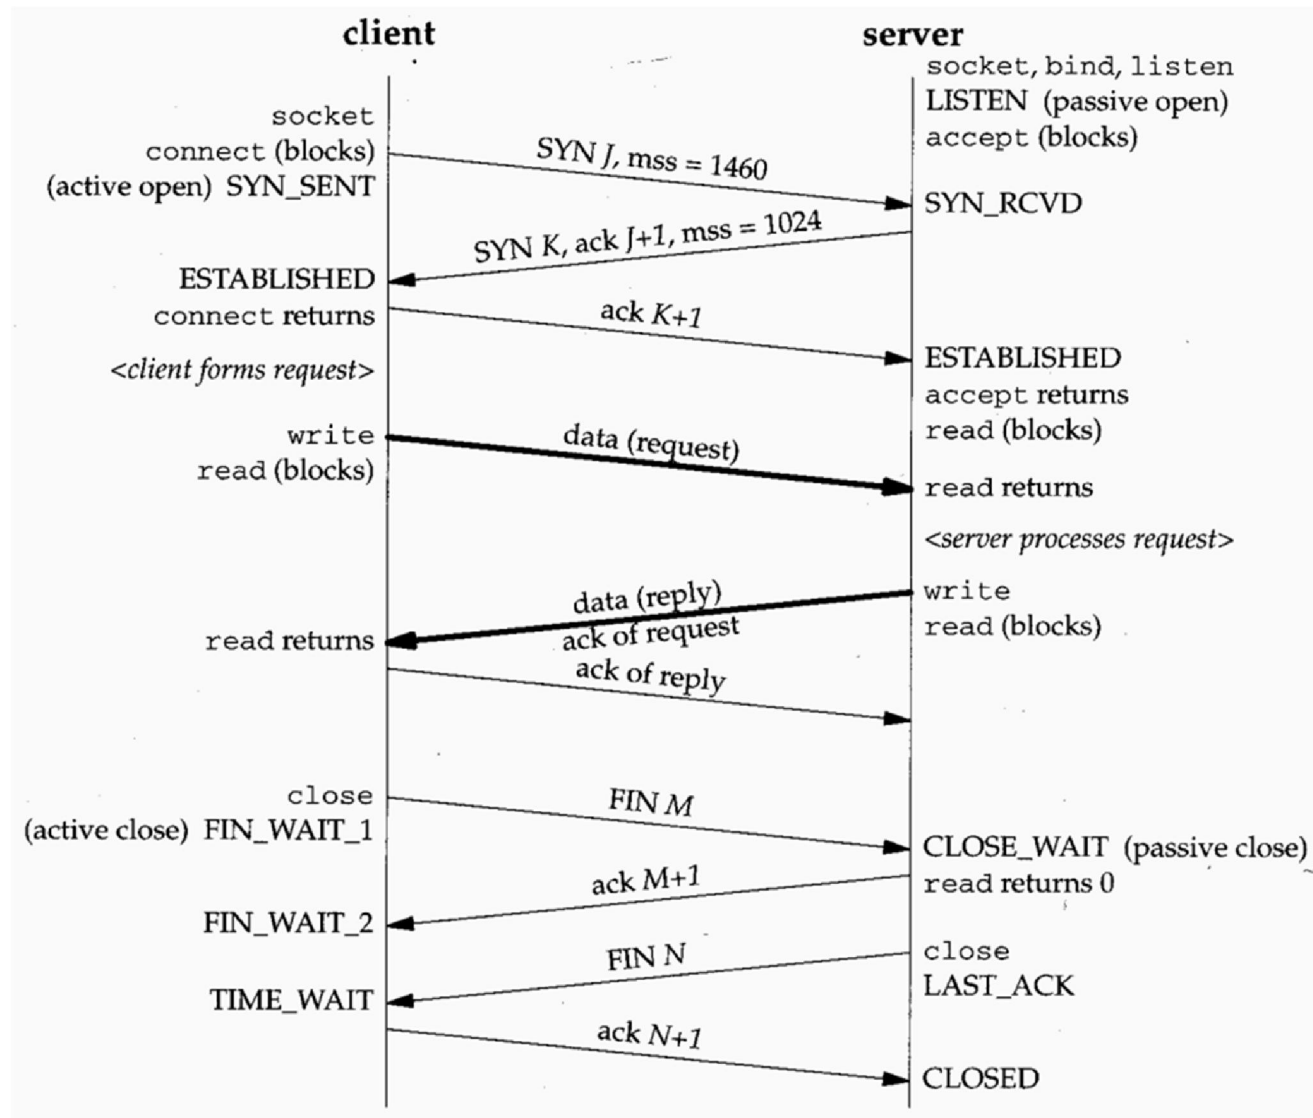
## Example *Closing* using a *Three-way Handshake*



## Example *Closing* using a Four-way Handshake (a Half Close)



# Example of a complete connection sequence



# TCP State Transition Diagram - explained

---

- ◆ The TCP ***State Transition Diagram*** shows the operation of TCP's connection *establishment* and *termination* phases:
  - There are 11 different states defined for a connection
  - Client and server transitions are shown as dark solid and dark dashed lines respectively
  - The transition from one state to another depends on the segment received in that state
    - e.g. an application performing an *active* open in the CLOSED state moves to the SYN\_SENT state and then to the ESTABLISHED state upon receipt of a SYN with an ACK
- ◆ The ESTABLISHED state is where most data transfer occurs

# TCP State Transition Diagram - explained

---

- ◆ Termination of a connection:
  - From the ESTABLISHED state there are two possible transitions:
    - An application calling *close* i.e. an *active* close moves to the FIN\_WAIT\_1 state
    - An application receiving a FIN (i.e. a *passive* close) moves to the CLOSE\_WAIT state
- ◆ In rare circumstances it is possible for both ends to send SYN/FINs simultaneously (known as *simultaneous open/close*):
  - This scenario is not explored here

# The TCP State Transition Diagram

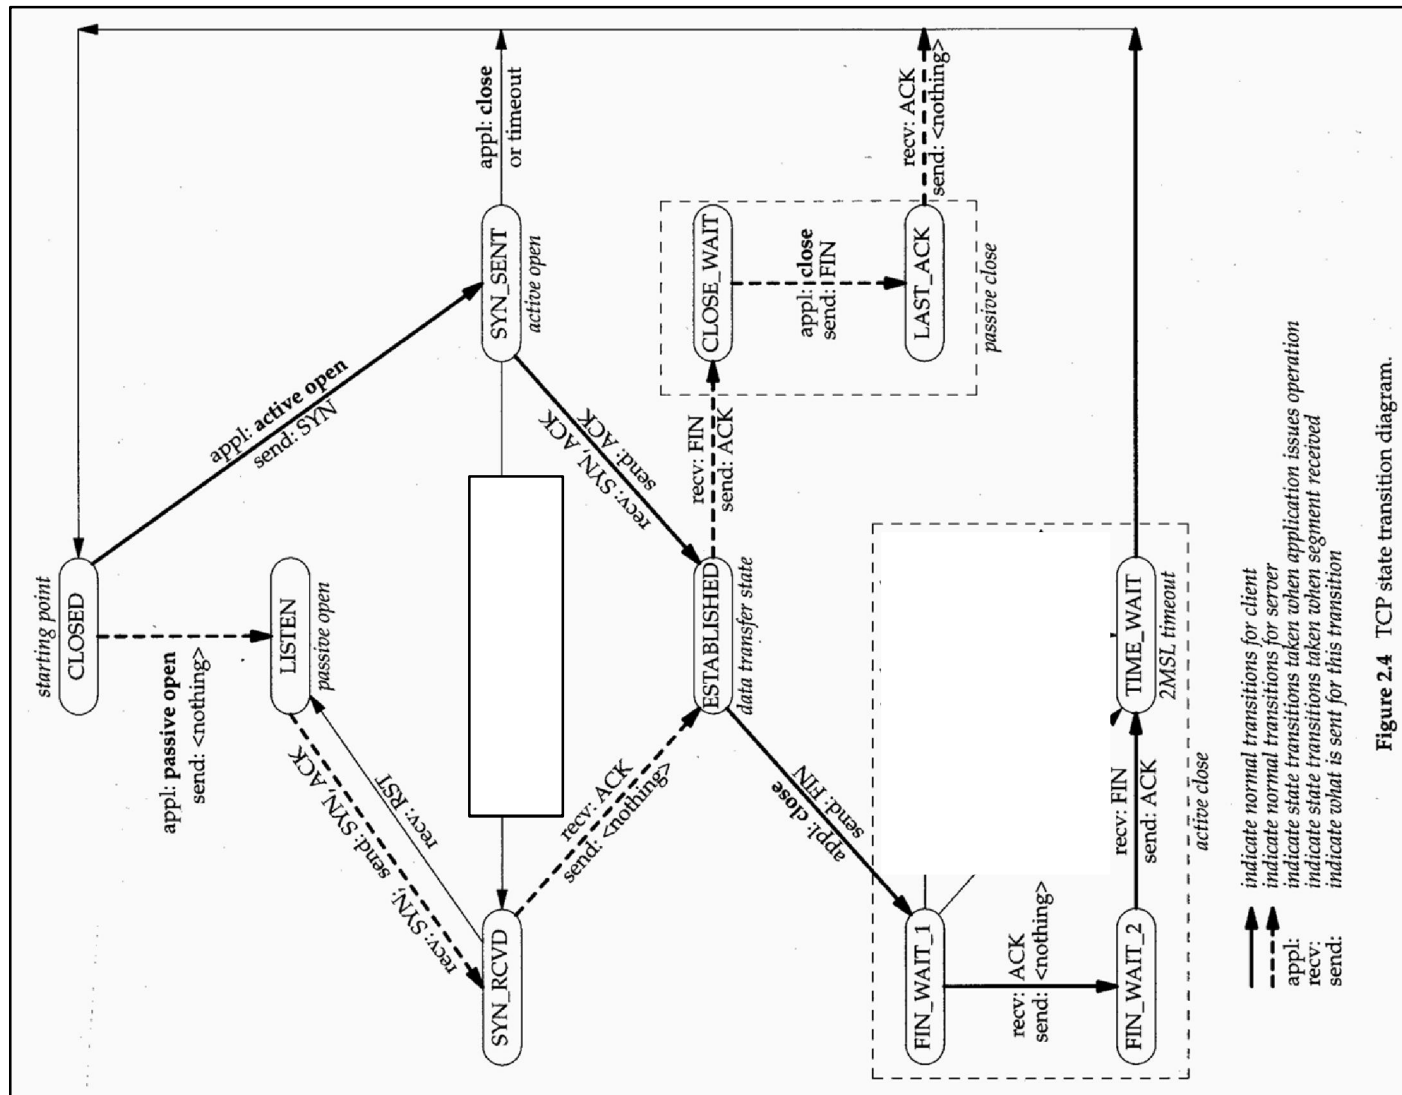


Figure 2.4 TCP state transition diagram.



# TCP's TIME\_WAIT State - explained

---

- ◆ The end that performs the *active* close goes through the TIME\_WAIT state for a period of twice the MSL (maximum segment lifetime) a.k.a. *2MSL*
  - MSL is the maximum amount of time that an IP *datagram* can live in an internet. This is linked to the TTL field (max value is 255) – to be covered later
  - TCP chooses a value for MSL of between 1 min. and 4 minutes

# TCP's TIME\_WAIT State - explained

---

- ◆ There are two reasons for the TIME\_WAIT state:
  - To implement TCP's full-duplex connection *termination* reliably
    - Recall two of TCP's service offerings is *full-duplex communication* and *reliable termination*
    - The end that performs an *active* close remains in the TIME\_WAIT state as it might have to retransmit the final ACK
  - To allow old *duplicate* segments to expire in the network:
    - Datagrams containing TCP segments can get caught in *routing loops* within an internet due to *routing errors*. These are known as *lost* or *wandering duplicates*
    - TCP must handle these duplicates for connections that have been *reincarnated*
    - TCP will not initiate an *incarnation* of a connection that is currently in the TIME\_WAIT state
    - This guarantees that all old *duplicates* from previous incarnations have expired