

Data Link Flow Control - Recap (Sliding Windows)


- ◆ This technique allows *multiple* frames to be in transit simultaneously
- ◆ Both stations use an extended buffer size to hold multiple frames
- ◆ The Sending/Receiving stations maintain a list of frames already sent/received
- ◆ This technique allows for much more *efficient link utilization*
- ◆ The transmission link is effectively treated as a *pipeline* that can be *filled* with many frames in transit *simultaneously*

Sliding Windows Flow Control


- ◆ Stations A and B each allocate buffer space for W frames
 - i.e. Station B can *accept* W frames and Station A can *send* W frames without any acknowledgement being sent or received
- ◆ Each frame contains a *sequence number*
- ◆ Station B sends *acknowledgements* that include the sequence number of the *next* frame expected
 - i.e. Station B is prepared to receive the next W frames starting at the *sequence number* indicated e.g. RR5

Sliding Windows Flow Control


Transmitters Perspective



Receiver's Perspective



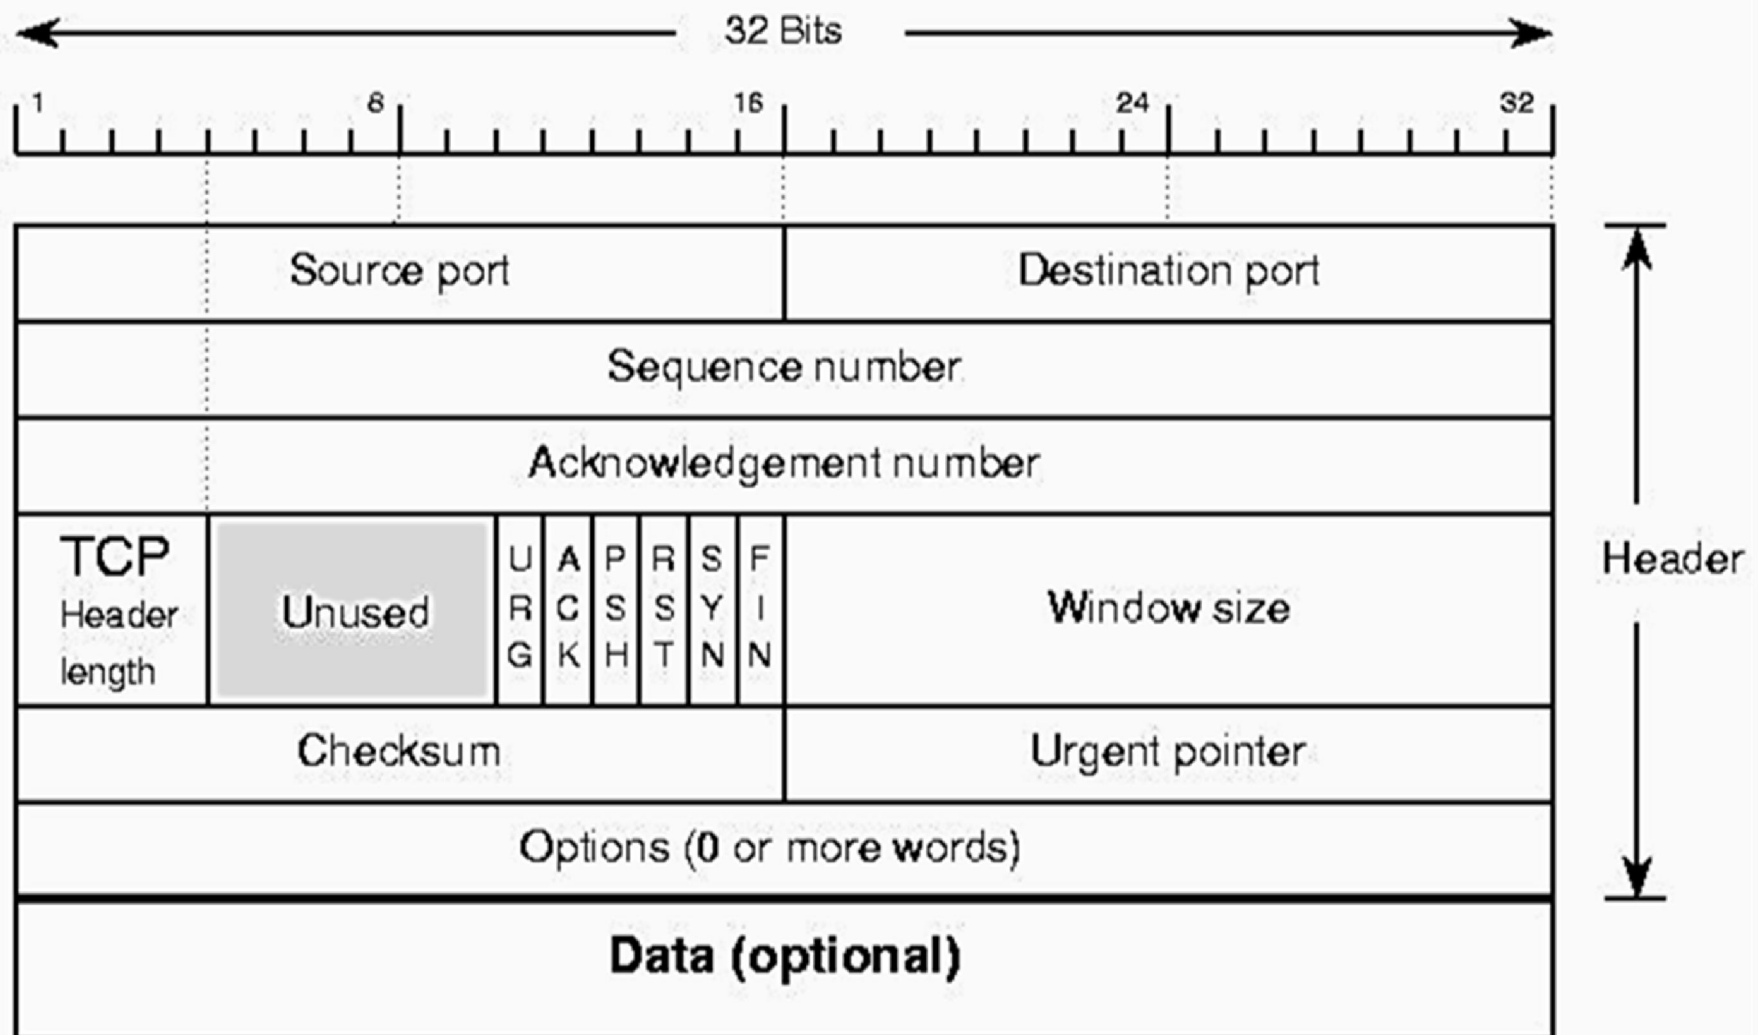
Example Sliding Windows



Sliding Windows Flow Control

- ◆ *Multiple* frames can be *acknowledged* using a single control message (*implicit acknowledgement*)
 - e.g. Receipt of ACK for frame **2** (RR3) followed later by ACK for frame **5** (RR6) *implies* acknowledgement of frames **3** and **4**
- ◆ Station A maintains a list of frame numbers it is allowed to *send*
- ◆ Station B maintains a list of frame numbers it is prepared to receive
- ◆ These lists can be considered as *windows*

TCP Segment Format




TCP Flow Control – *Buffers and Windows*

- ◆ Recall that TCP creates two buffers per socket:
 - One for incoming data (RECVQ)
 - One for outgoing data (SENDQ)
- ◆ The incoming buffers can easily overflow
- ◆ To prevent this the receiving TCP entity uses a ***window mechanism***

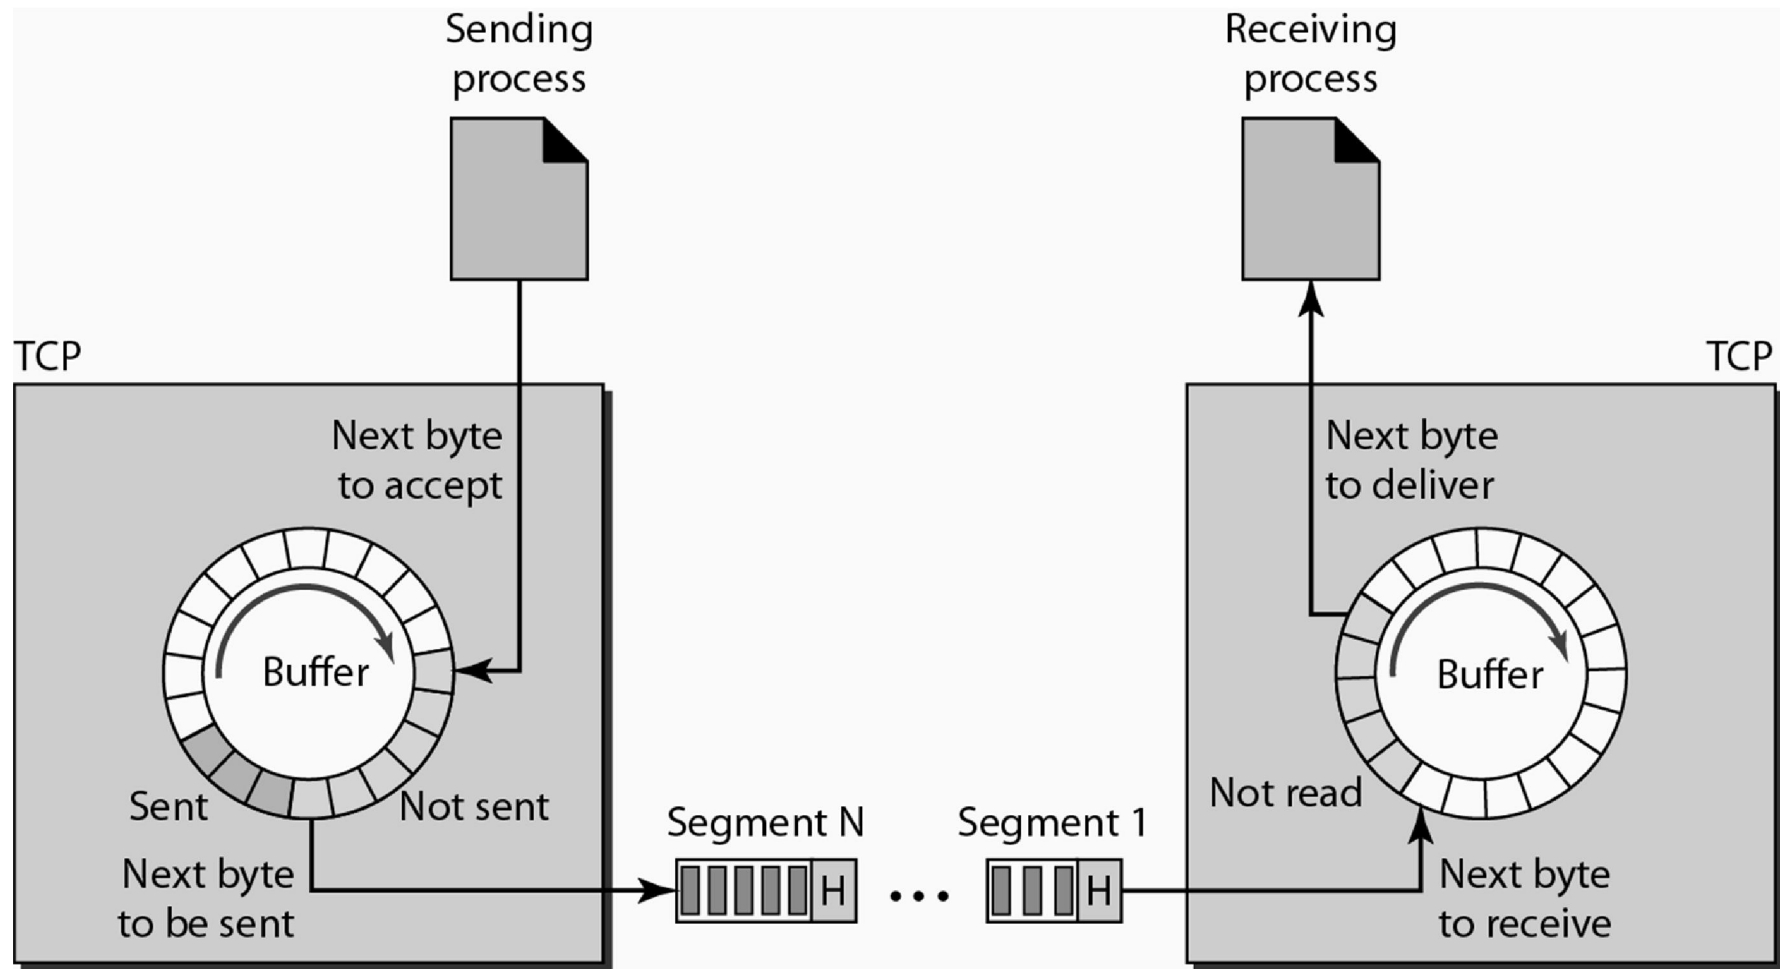
TCP Flow Control – *Buffers and Windows*

- ◆ Each end of the connection allocates a ***window*** to hold incoming data:
 - The size of the initial window is set during *Phase 1, Connection Initialization*, when both sides send their SYN messages (using the *Window Size* field)
 - Thereafter, throughout *Phase 2, Data Exchange*, all ***Ack*** messages include a ***window advertisement***
- ◆ The *window advertisement* can be positive or zero depending on space availability in RECVQ.

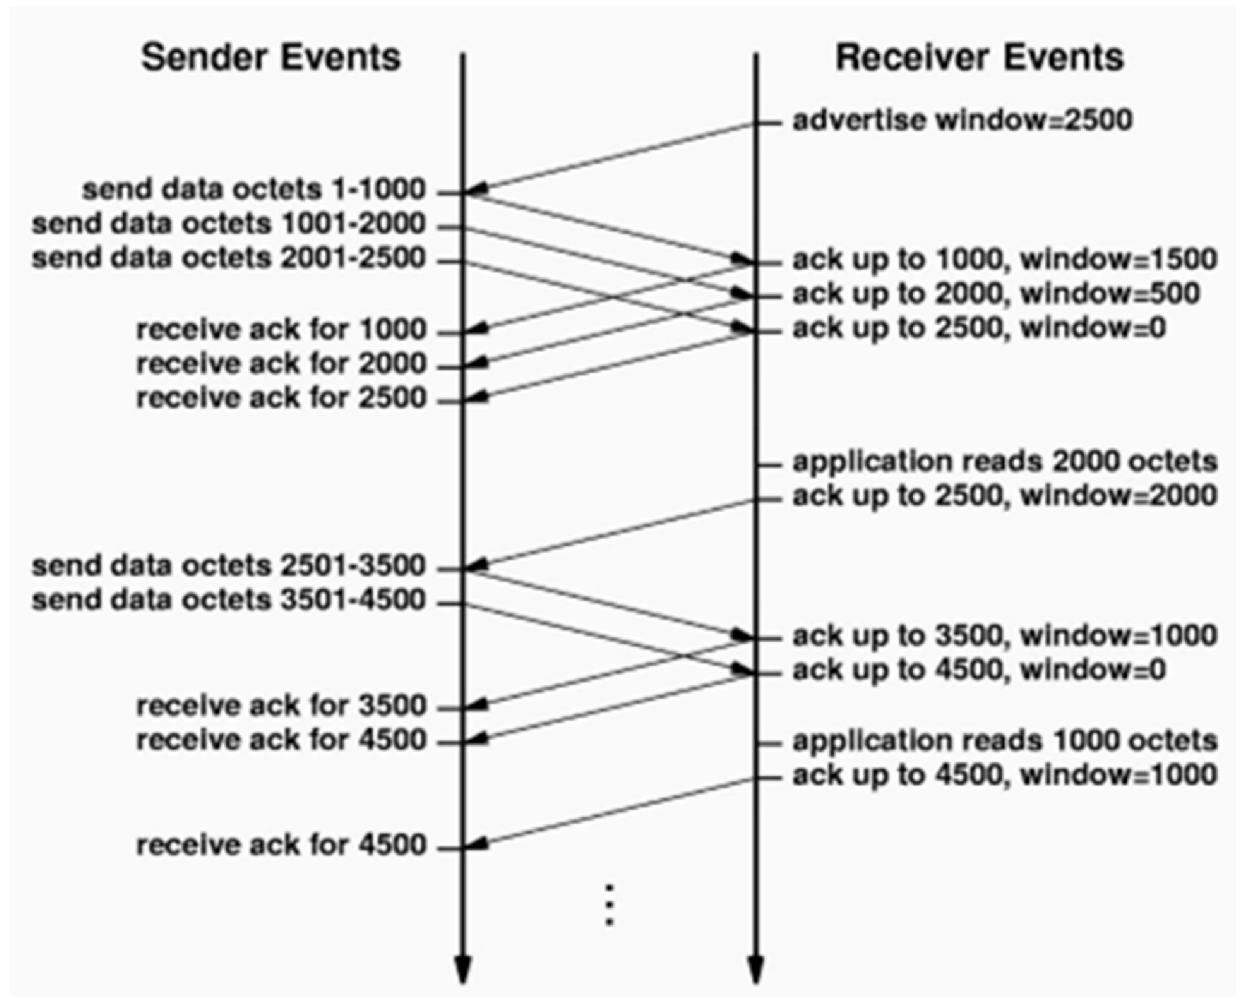
TCP Internal Data Buffers



TCP Internal Data Buffers



Operation of *window advertisements*



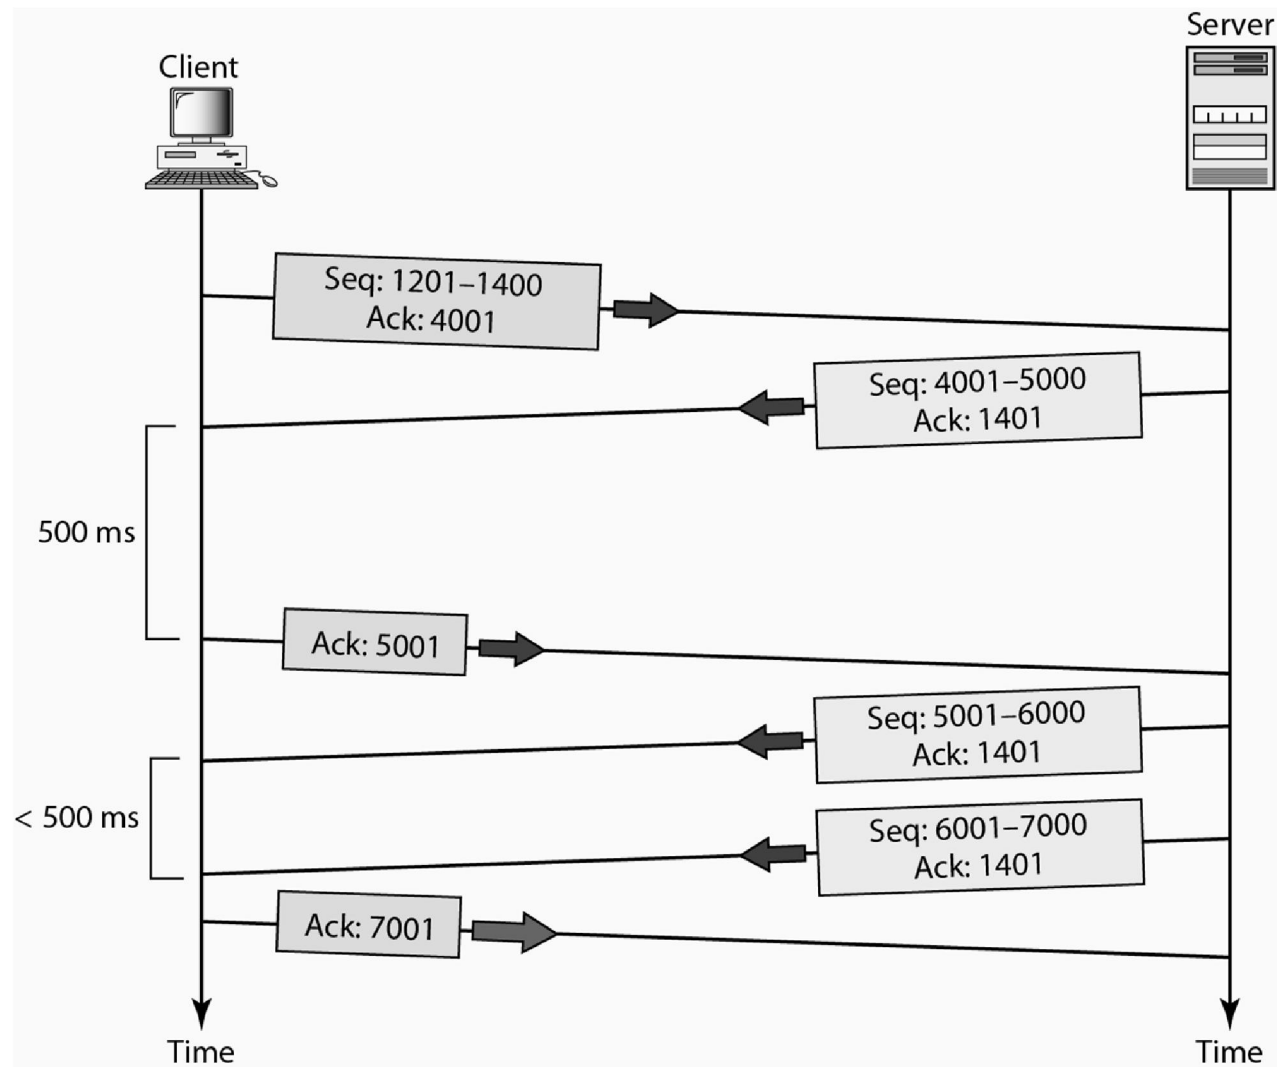
Error Control - Achieving Reliability

- ◆ TCP must address the following reliability problems:
 - Unreliable delivery by the underlying communication system
 - Segments can be lost, duplicated, delayed, or delivered out of order by the underlying communication system
 - Computer reboot
 - Segments from previous connections can arrive. The receiving host having just re-booted will not know what to do with them
- ◆ To overcome these problems TCP must:
 - Ensure that all messages are unambiguous
 - Be able to deal with *lost packets*

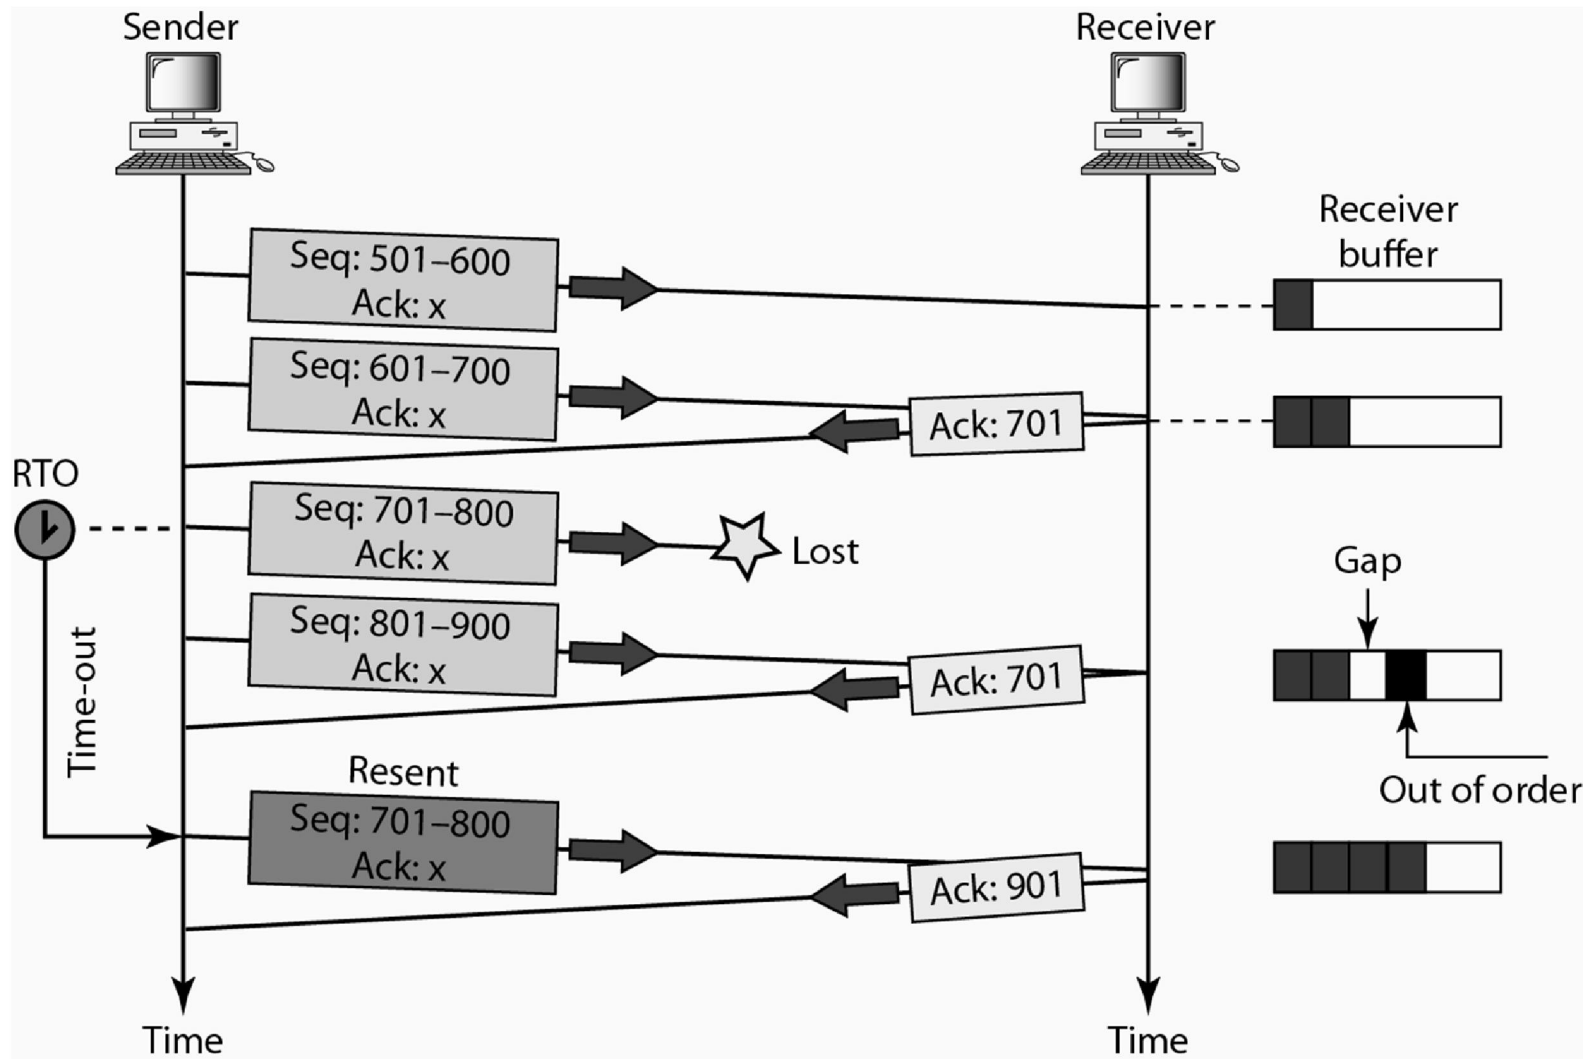
Handling Packet Loss - Retransmission

- ◆ TCP implements a *retransmission scheme*
- ◆ This involves sending *acknowledgements* and the use of *timers*
- ◆ This scheme is the key to the success of TCP
- ◆ How long should TCP wait before retransmitting?
 - It depends upon:
 - The underlying network
 - Traffic levels
- ◆ TCP uses an *adaptive retransmission* scheme

TCP Data Flow – Normal Operation



TCP Data Flow – Lost Segment Scenario



TCP's *Adaptive Retransmission* Scheme

- ◆ TCP estimates the *round-trip delay* for each active connection
- ◆ Using a *weighted average* together with a *variance* factor TCP calculates a value for the *retransmission time*
- ◆ This helps TCP to react quickly to changes in *traffic levels* and to maximize throughput on each connection